

GS 1380

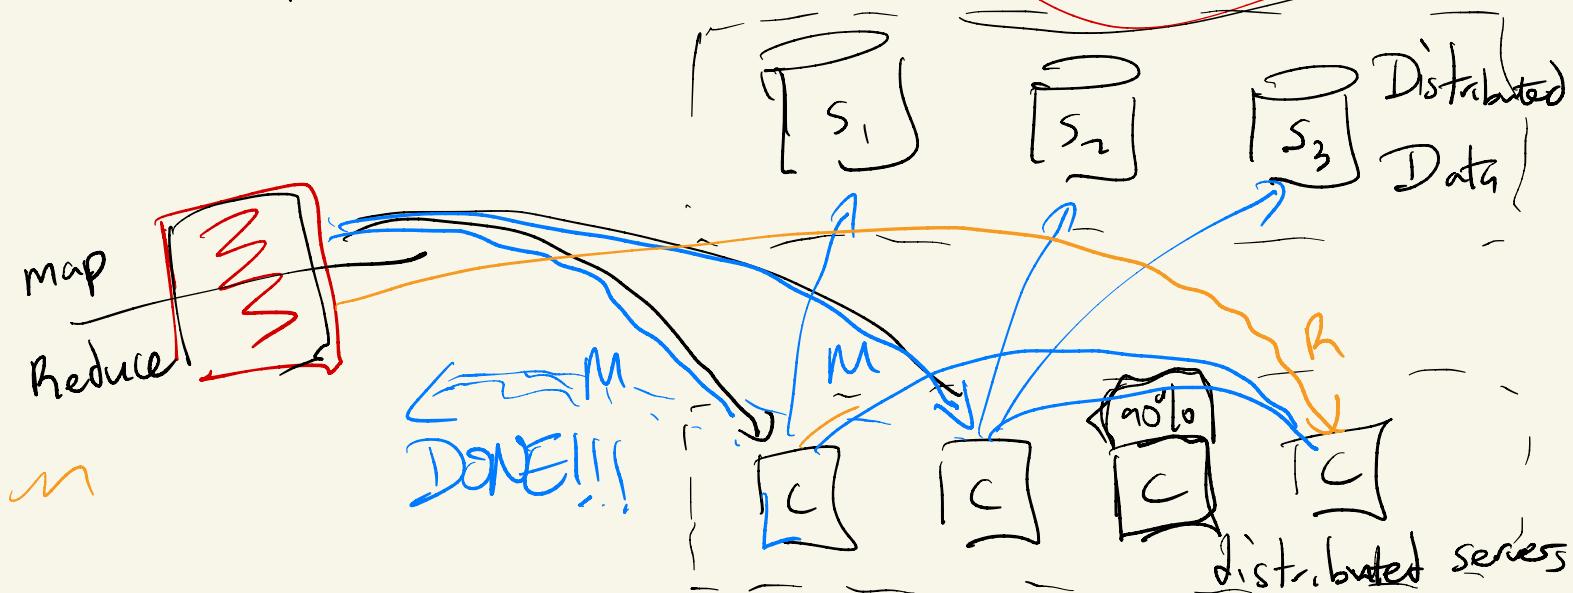
Day 2 : Jan 26

Theophilus Benson



- ① Motivation for D.S. concepts
(why is D.S. Hard)
- ② Networking (How do this enable D.S.)
- ③ Failures (types | practicality | overheads)
- ④ "fallacies of distributed computing"

Map Reduce = large scale + easily distribute workloads + Massive data



- ① pick servers for Mapping
- ② for each mapper inform it of storage to use

m

- ③ pick servers(s) for reduce
- ④ inform mappers of reduce location

- Requirements
- a) keep track of all servers
 - (i) Health
 - (ii) progress
 - (iii) data location
 - (iv) load utilization

Count (*) from

all data

where name = 'theo'
& age ≥ 12

||
reduce

||
map

> 2%

why is the master node a bad idea?

- ① single point of failure
- ② all load is one server

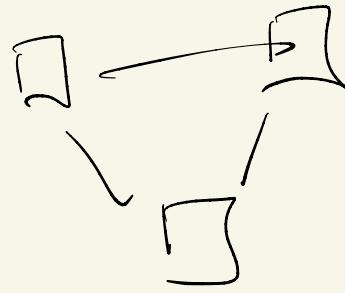
- ① cheap
- ② easy to write
- ③ easy to coordinate
- ④ how much up time is good enough?

Does simplicity still matter with formal verification?

failure / performance problems

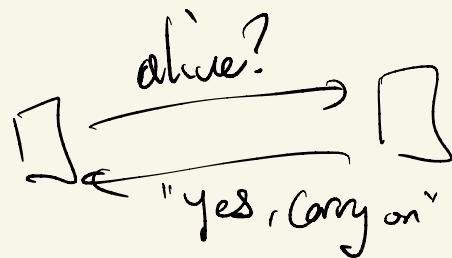
① failure types

- Ⓐ network
- Ⓑ Power
- Ⓒ resource exhaustion (CPU/Mem)
- Ⓓ restart (state failure)
- Ⓔ Hardware failure



easy to detect

- ① Server / power
- ② Some n/w failures

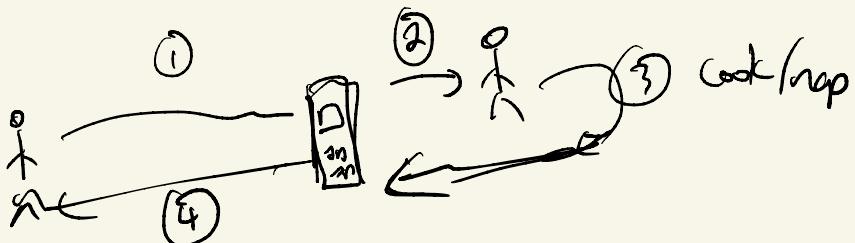
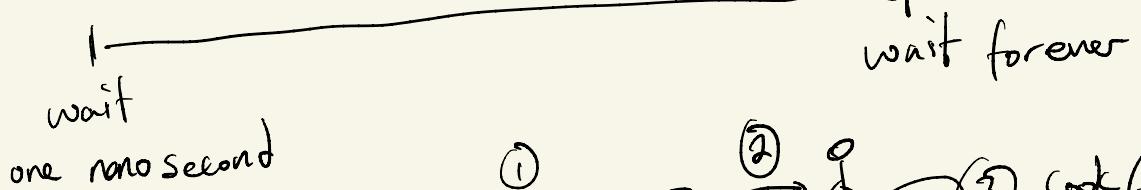


failure detection = heartbeat messages

Questions about designing

- ① frequency
- ② performance impact
- ③ msg size
- ④ much data to detect failure

(1) How long should you wait for a reply?



$$\frac{1s}{1} + \frac{1hr}{2} + \frac{3.5hr}{3} + \frac{1s}{4}$$

① past experience

② 12 hr timezone

③ day / versus night

④ indication of receipt

(2) How often should you send heart beat msgs?

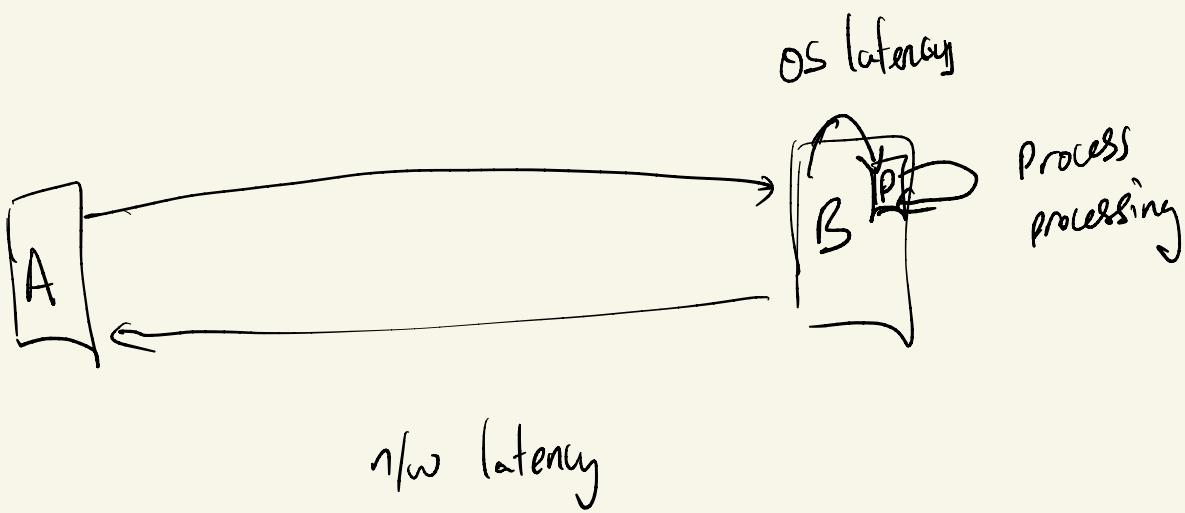
every second

every two years

(1) depends on CPU (speed/generation)!!

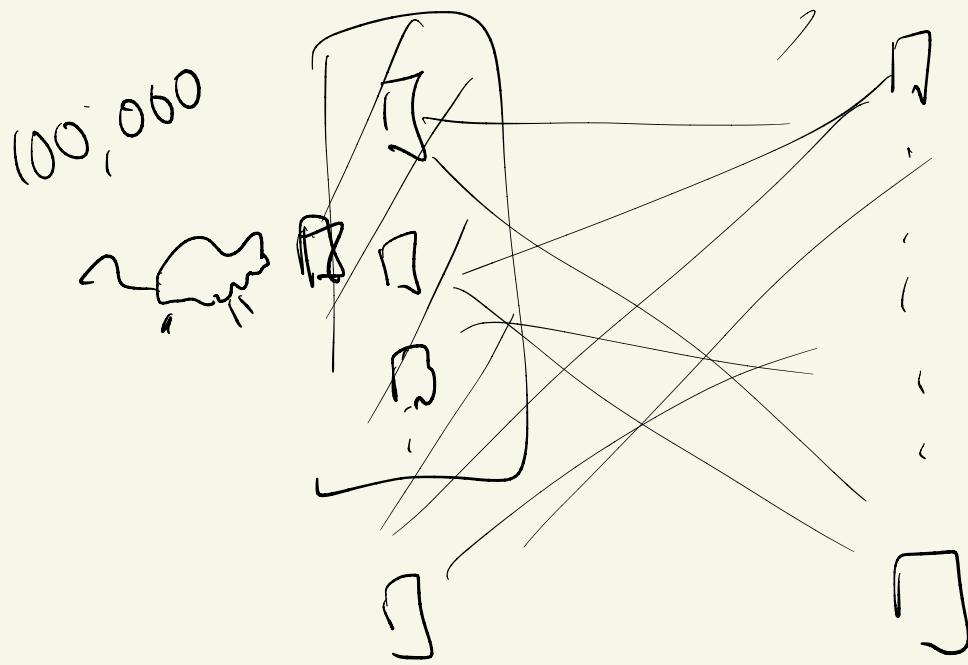
(2) criticality of functionality!

Backup ↘ ↗ Google search



$$\text{Heart beat} \geq n/w + \text{OS} + \text{process} + n/w$$

Heart beat time is too small → mistakes (where you think the server is down but it's up)



Theo. Com

- ① 100 k servers
- ② power outage takes out 30 k
- ③ after outage all servers are online