

CSCI 1380

Day 3 [Jan 28]

Theophilus Benson



Last Class

- (1) Motivation for D.S. (i.e. MapReduce)
- (2) Failures (heartbeats)

This Class

- (3) Networking
- (4) Fallacies of distributed computing
- (5) Performance optimizations

First: the heartbeat dilemma

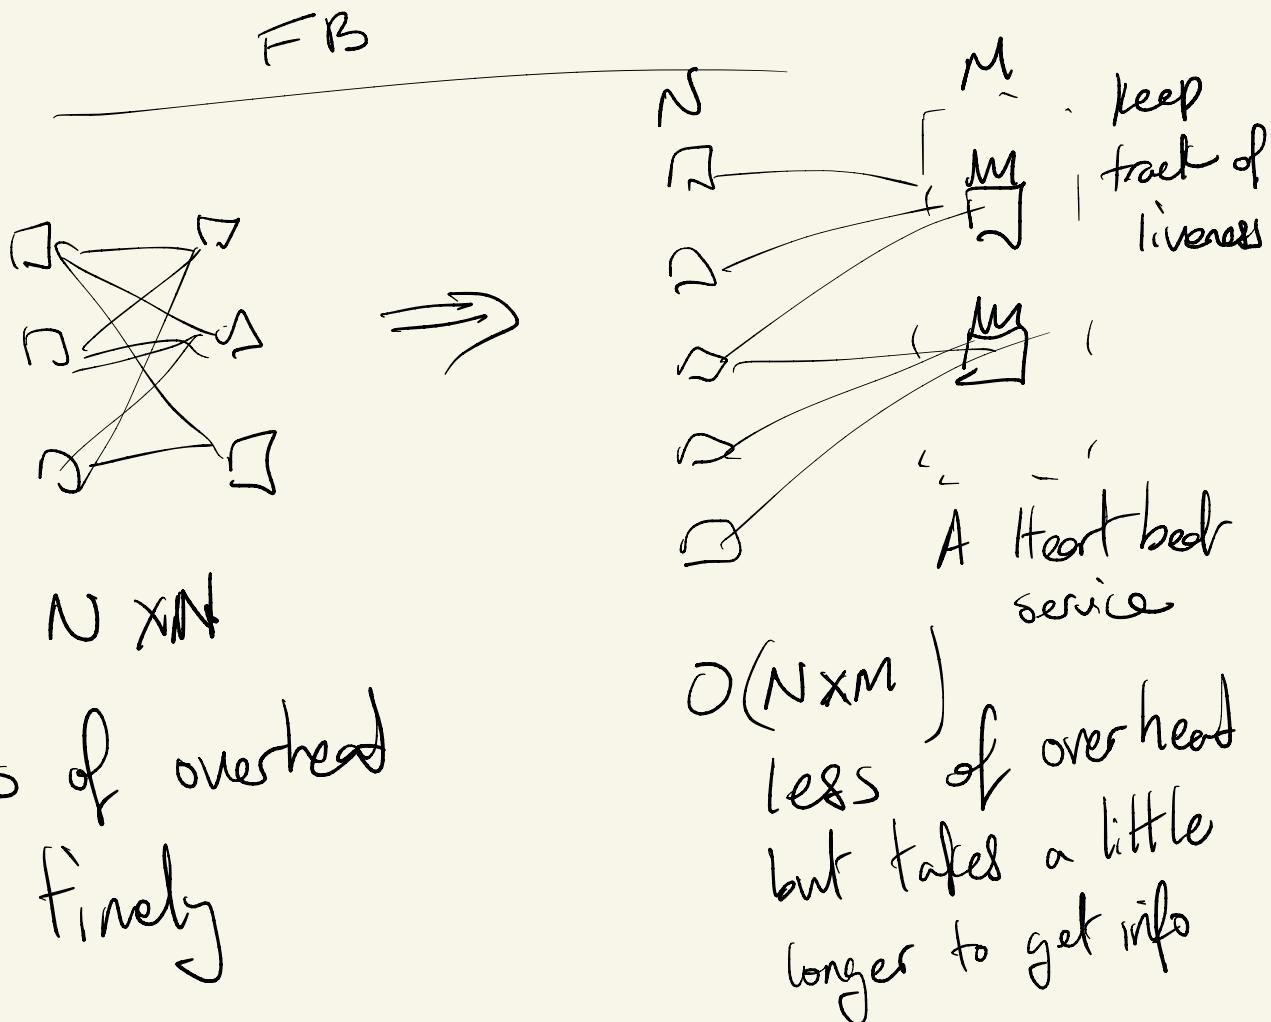


a company with
100k servers.

30k of these servers
fail & restart. What
sort of bad things
can happen?

AWS

new cluster of machines are added
all machines try to connect to each other
a thread to manage connections
solution : more hardware



Protocols (TCP / UDP)

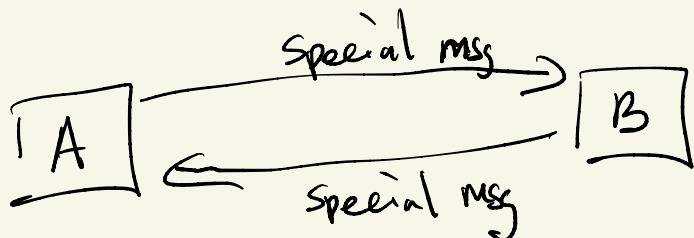
TCP

- ① need to make a connection with special msg (TCP handshake)

- ② after handshake: send/receive data

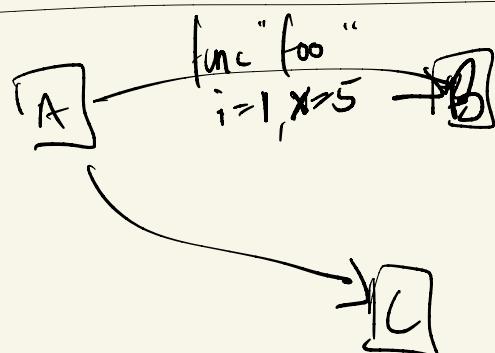
- ③ special rate control algorithm & retransmit algorithm

You never worry about lost packets & TCP lets you share the network fairly



Problems

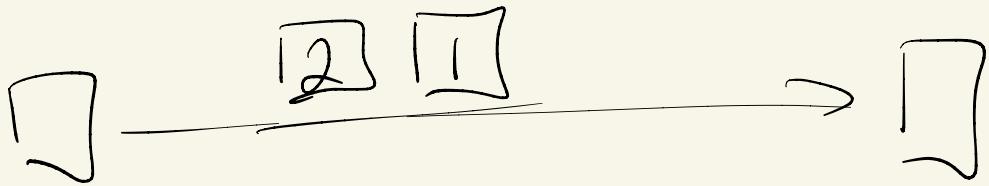
- ① overheads
② lots of messages
③ no one cares about lost packets



remote function calls \Rightarrow UDP
(RPC)

You want to avoid the startup overhead of TCP

UDP (reordering) || loss)



How to go from monolithic to distributed ?

App - c

- ① makes distributed calls : needs info about other components

② should we create each as a unique process ?

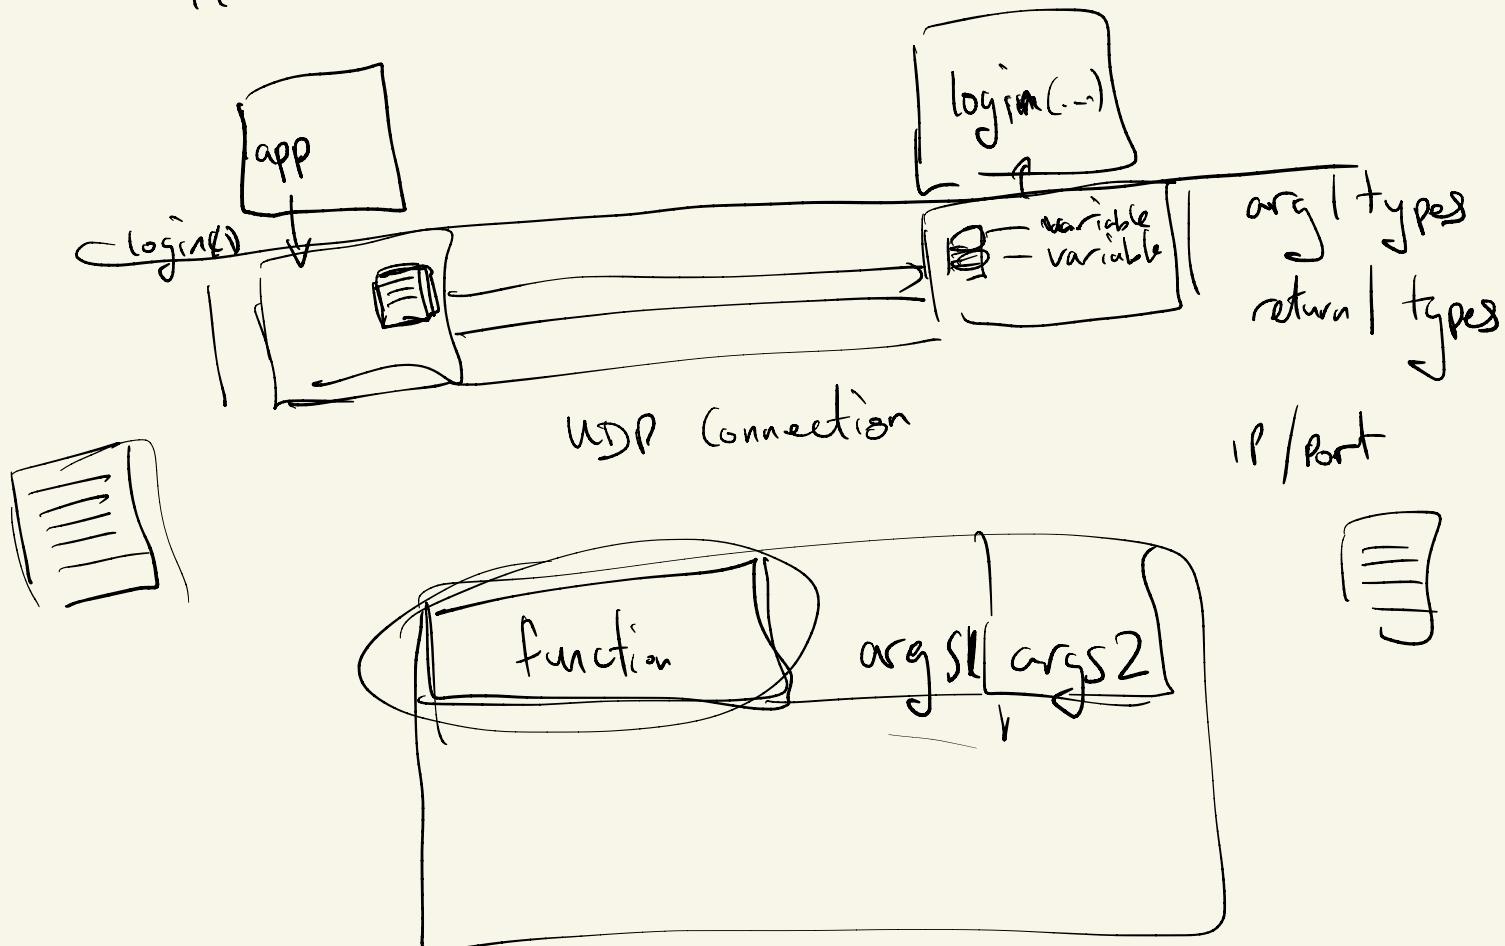
keep track of connections

metadata : where are the processes located ? "Port"

IP address + port

- ① each function becomes a process
- ② "app" needs information about the other process
- ③ "Special" to make network connections

RPC \Rightarrow Remote Procedure Calls



- (1) Heartbeats (when things go wrong)
 - (2) N/W \Rightarrow (TCP v. UDP)
 - (3) Socket code | RPC
-

Next class

RPC semantics | RPC structure

Performance | load balancers