

---

---

---

---

---

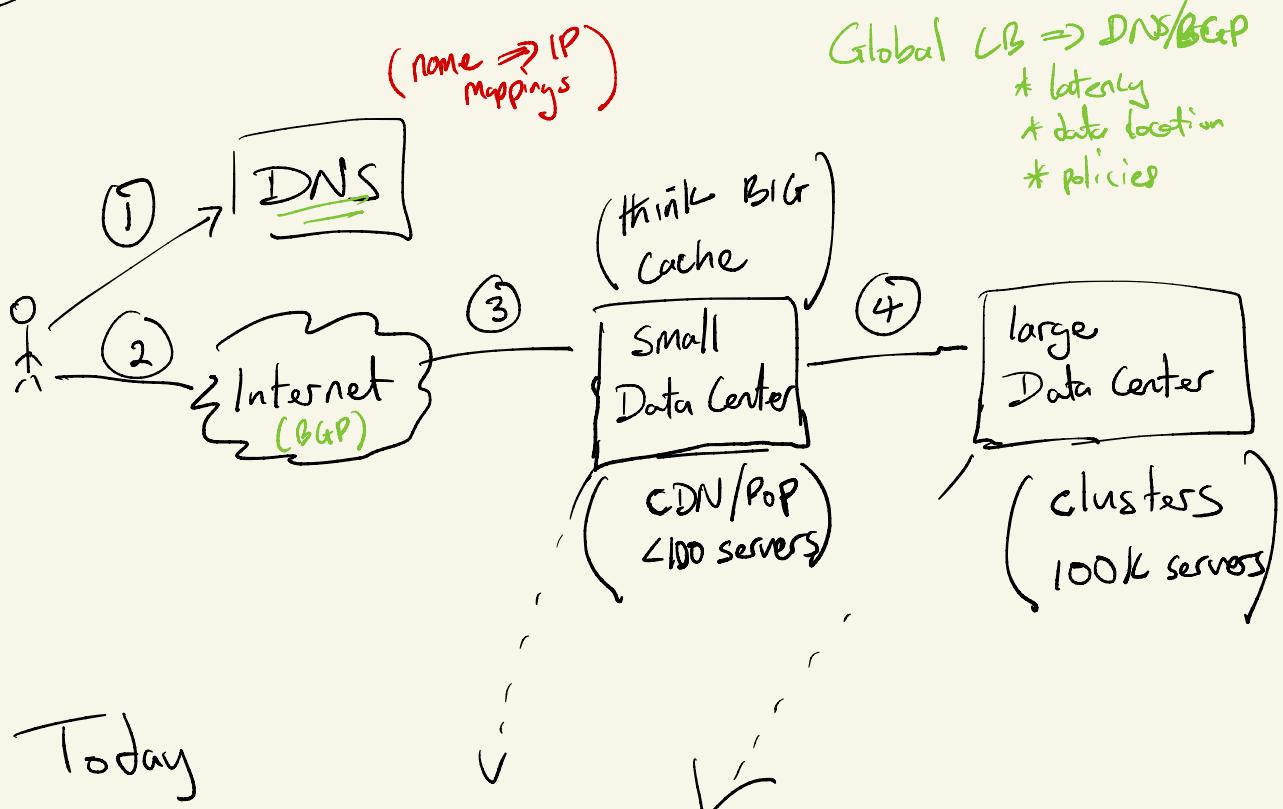


To day

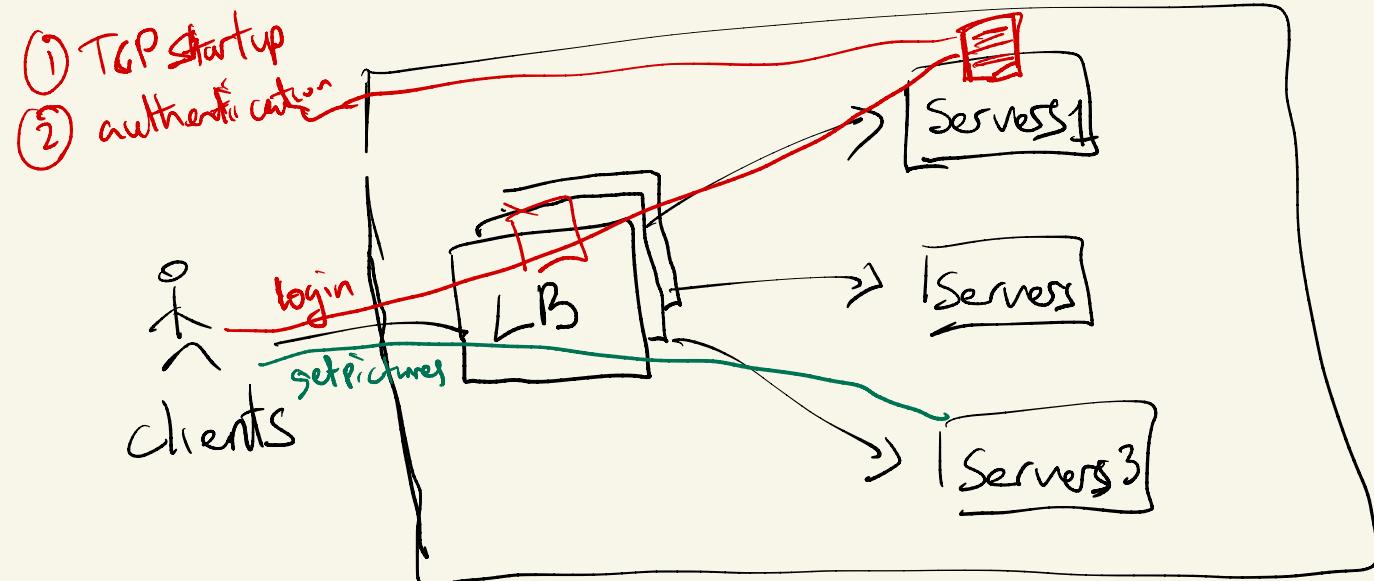
① Review of Day 6 (Please answer Top Hat Questions)

② Modulo L/B

③ Consistent hashing

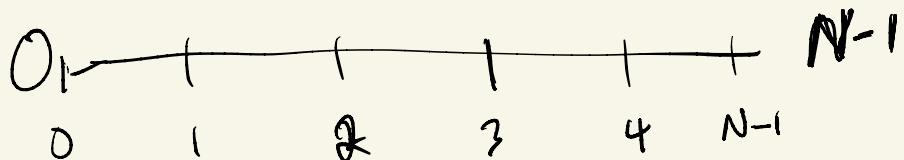
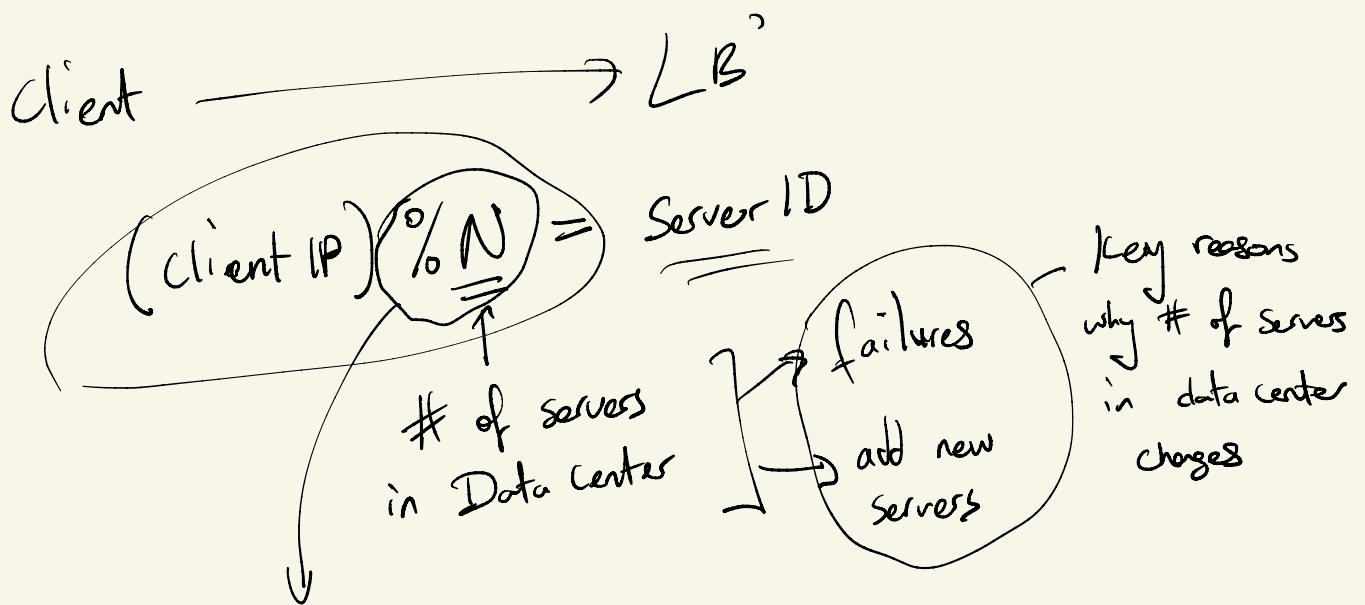


Today



$LB$ techniques	eBay's Neutrino	Google's Maglev	FB's Icaran	AWS ELB	GitHub LB	fastly
least connection	✓					
round robin	✓					
Chash (consistent hashing)	✓					
RHash (rendezvous hashing)		✓	✓			
Modulo LB	✓				✓	

## Modulo LB

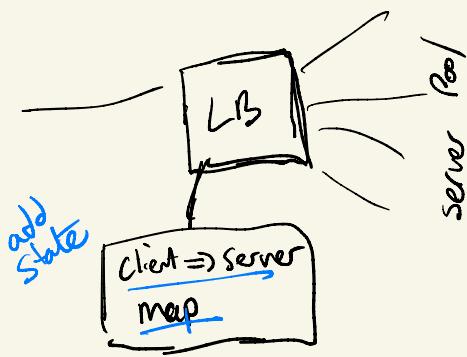


$$N: 6 \Rightarrow 7 \quad \left( \frac{28}{42} \right)$$

How can you extend LB (Random / Modulo) to provide consistent mapping of client to servers?

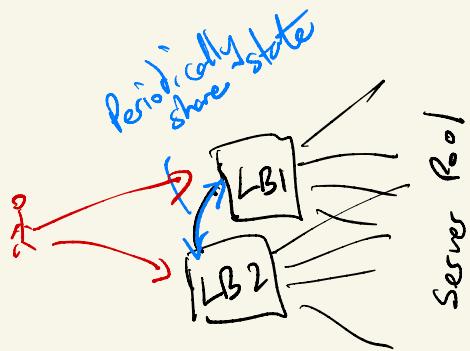
- ① LB should hold cache of client  $\Rightarrow$  server mapping
- ② store the mapping in an external location (disk or S3)
- ③ Random w/ Cache

## Extending Round Robin



on LB failure  $\Rightarrow$  find way to replicate state  
(external  $\Rightarrow$  S3 or persistent)

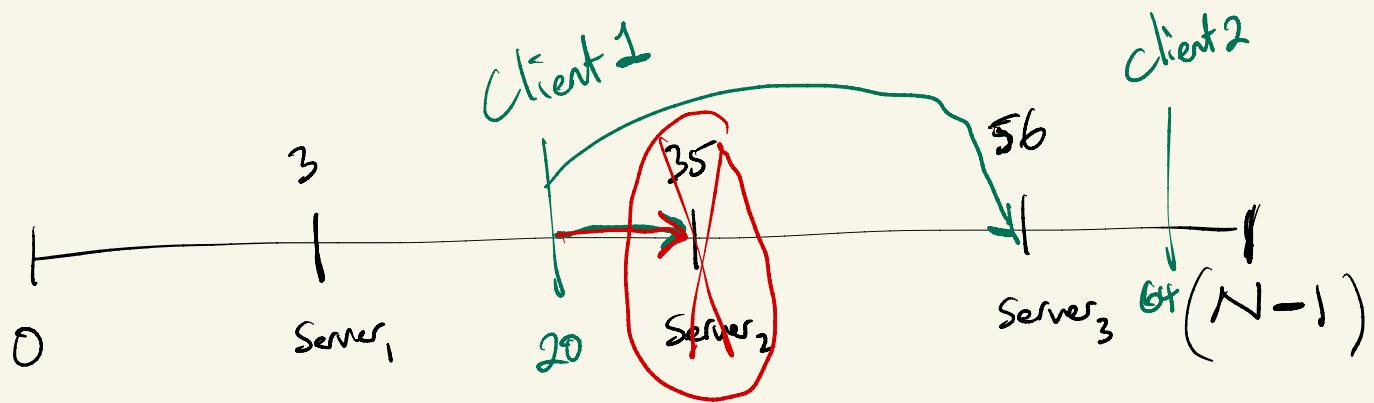
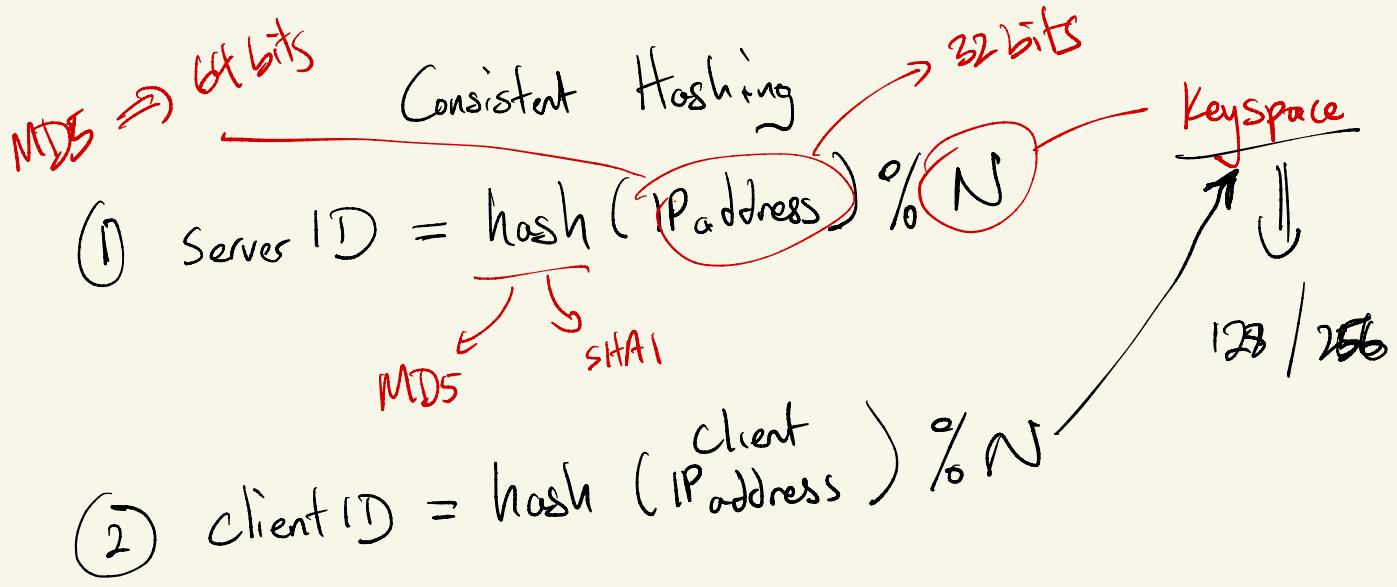
if many LB  $\Rightarrow$  find way to share state



- ① external k/v
- ② RPL calls to exchange state

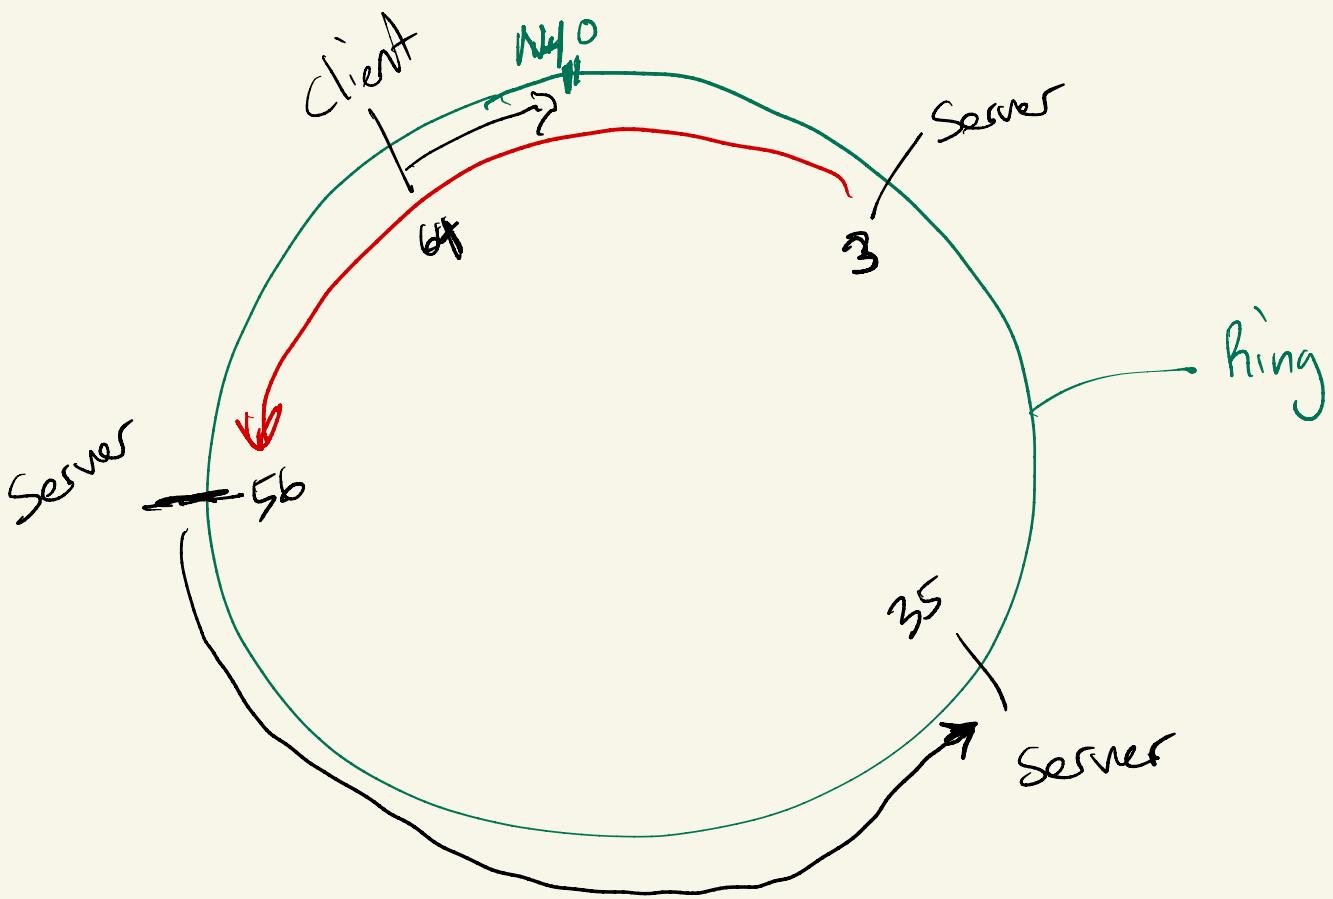
Maintaining & sharing state is complex  
introduces issues during LB failure/restart  
hard to share state fast enough

How can you be stateless & provide consistent LB?



using the equations  $\Rightarrow$  every LB can calculate server ID without state

- ① How to deal w/ hash collisions?
  - ⓐ Client  $\Rightarrow$  this okay: the clients are assigned to the same server
  - ⓑ Server  $\Rightarrow$  change IP address of the server



serverlist = sort server ID

foreach server in serverlist {

    if (server > clientID) {  
        myServer = server  
        break

} find first  
server with ID  
greater than  
client



$\rightarrow N$

④ Key space "wrap around" =

} if (myServer == nil)

    myServer = serverlist[0]

if no server w/ ID  
greater than client then you are in wrap around edge case  
then pick first server (which turns out to be  
next server when in  
"ring")

