

CSCI 1380

Day 8

Consistent Hash + Tapestry



Last Week

* Consistent hashing

Motivation: We need to consistently assign clients to servers
servers maintain client specific state (TCP/session)

Solution: Create serverID & keyID using some equation

$$ID = \text{hash}(\text{IP}) \bmod N$$

based # of servers / client

- Assumptions:
- (1) ServerID/KeyID are evenly distributed
 - (2) Load Balancer maintains a list of all servers

func



loadBalance (clientIP)

$$\text{ClientID} = \text{hash}(\text{clientIP}) \bmod N$$

Serverlist = sort serverIDs

need to know all servers

```
for each serverID in Serverlist {  
    if (serverID > ClientID)  
        return serverID
```

```
}  
return Serverlist[0]
```

wrap around case

$$\text{ServerIDs} = [3, 35, 56]$$

$$N = 64 \rightarrow [0-63]$$

Why can a system have uneven distribution of load?

- (1) Server failures (many servers failing)
- (2) uneven # of requests from clients
- (*) (3) Small # of servers which get hashed to similar location
- (4) Servers with different H/W

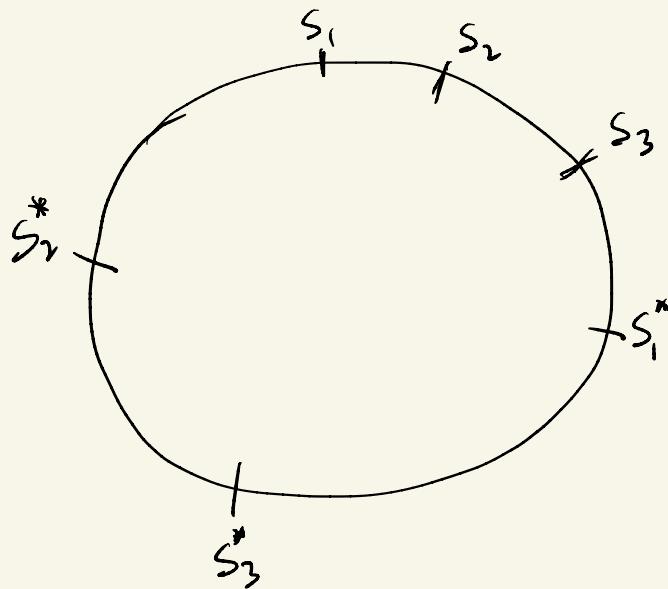
FB / Amazon

failures are temp

Get
more servers
Solution

(virtual nodes)
Virtual Server

(solution to
uneven distribution)



Virtual = give each server \overbrace{X} IDs

- ① have X hash functions
- * ② give each server X IPs
- * ③ "salting" the IDs

$s_1 \leftarrow 10.10.10.100$
 $s_1 \leftarrow 10.10.10.10$
 $\text{hash}(xxx\ 10.10.10.10) \leftarrow \text{salting}$

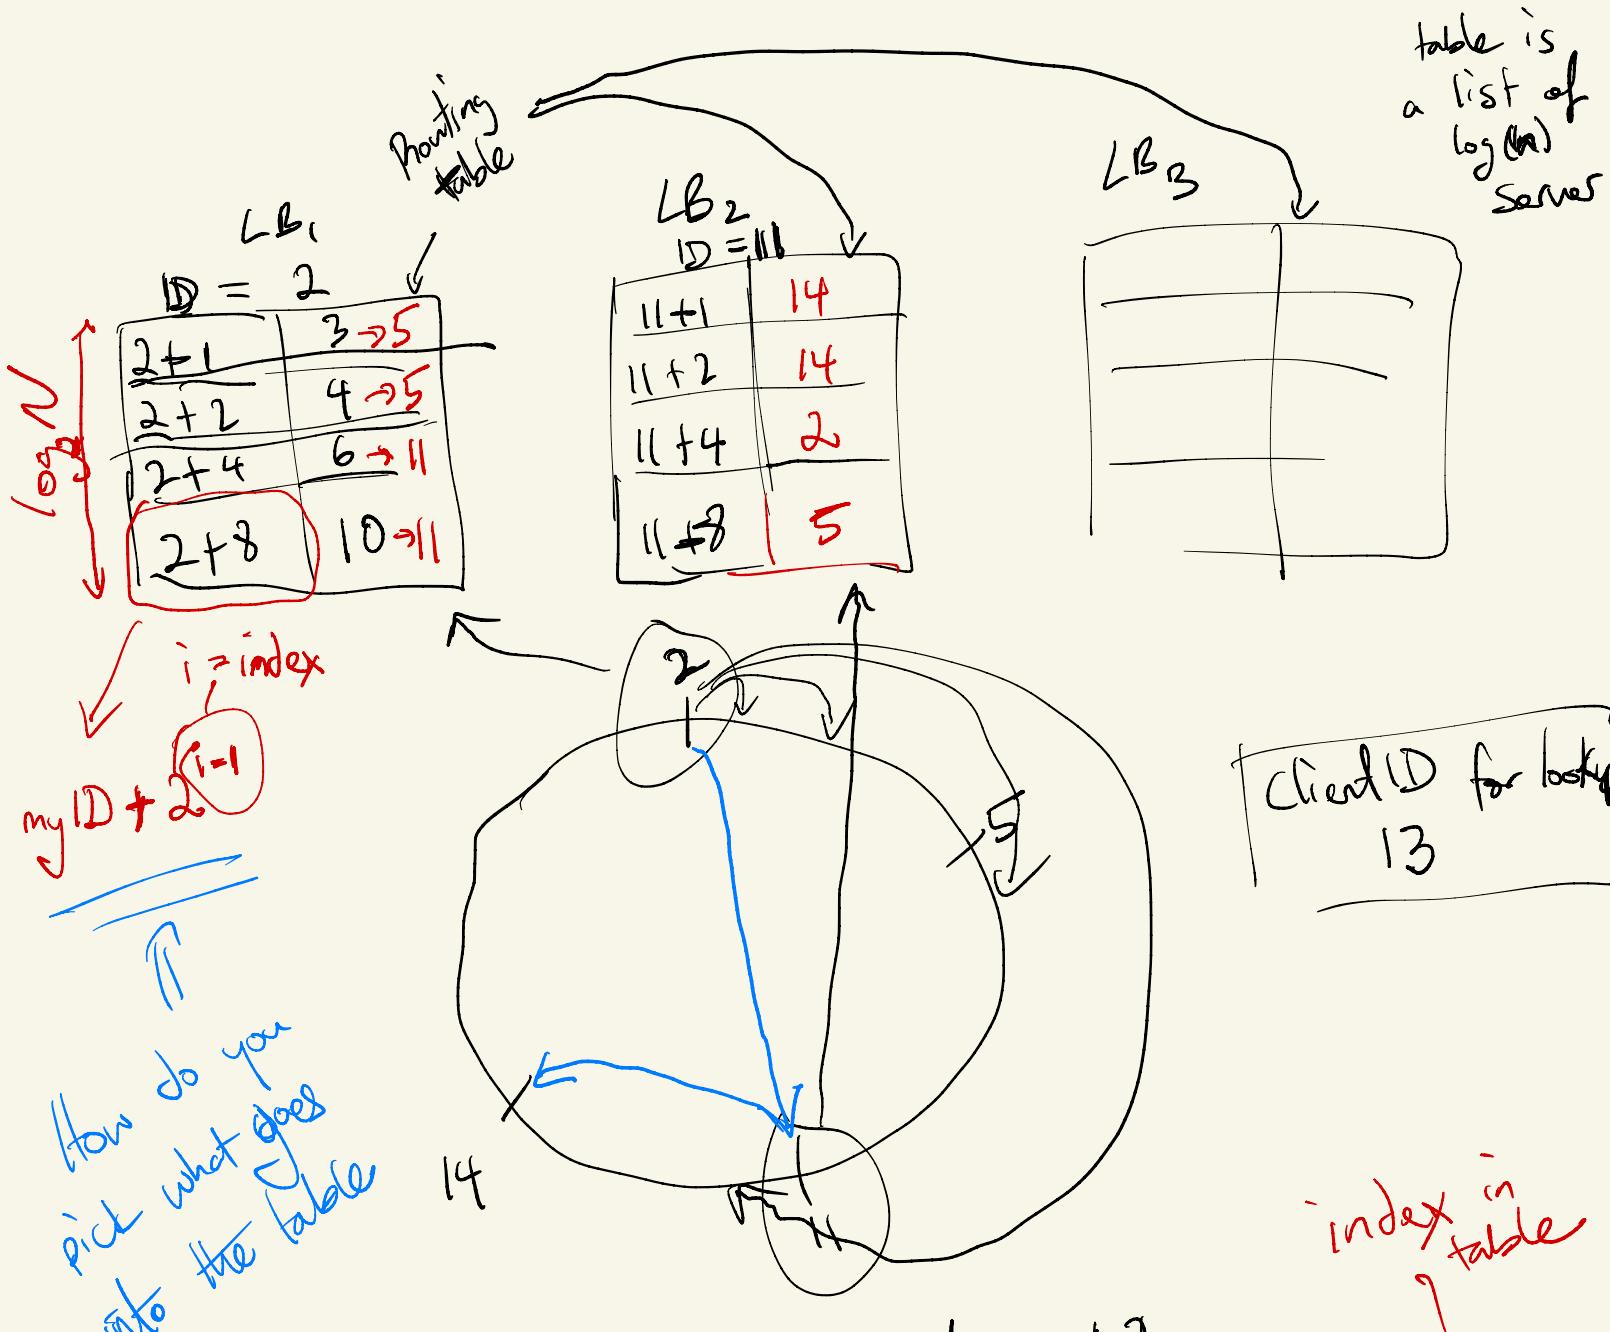
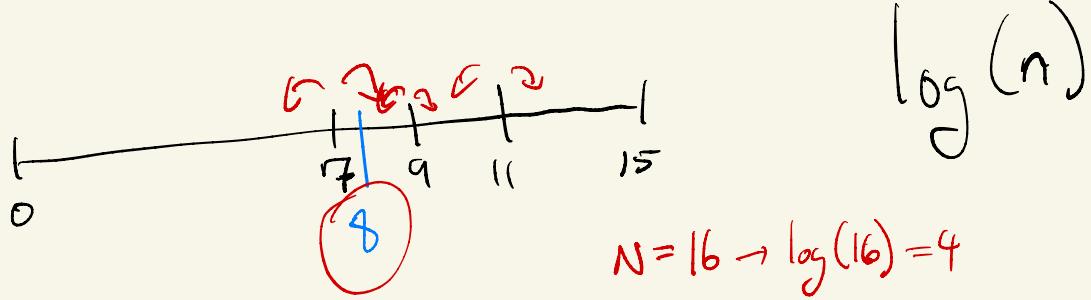
Hash ($xxx + IP$) mod N

append random string

How do you set " X "?

minimize the size of server IDs that
each LB needs to maintain?

How does binary search work?



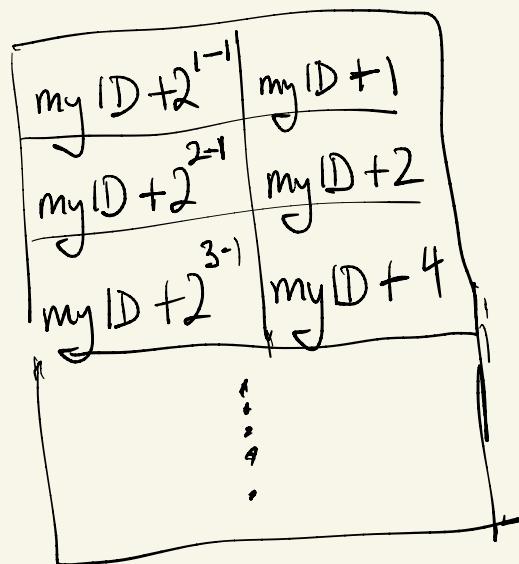
lookup for 13 starting at servedID

Each entry in table is calculated: $\text{myID} + 2 - i$

Chord recap

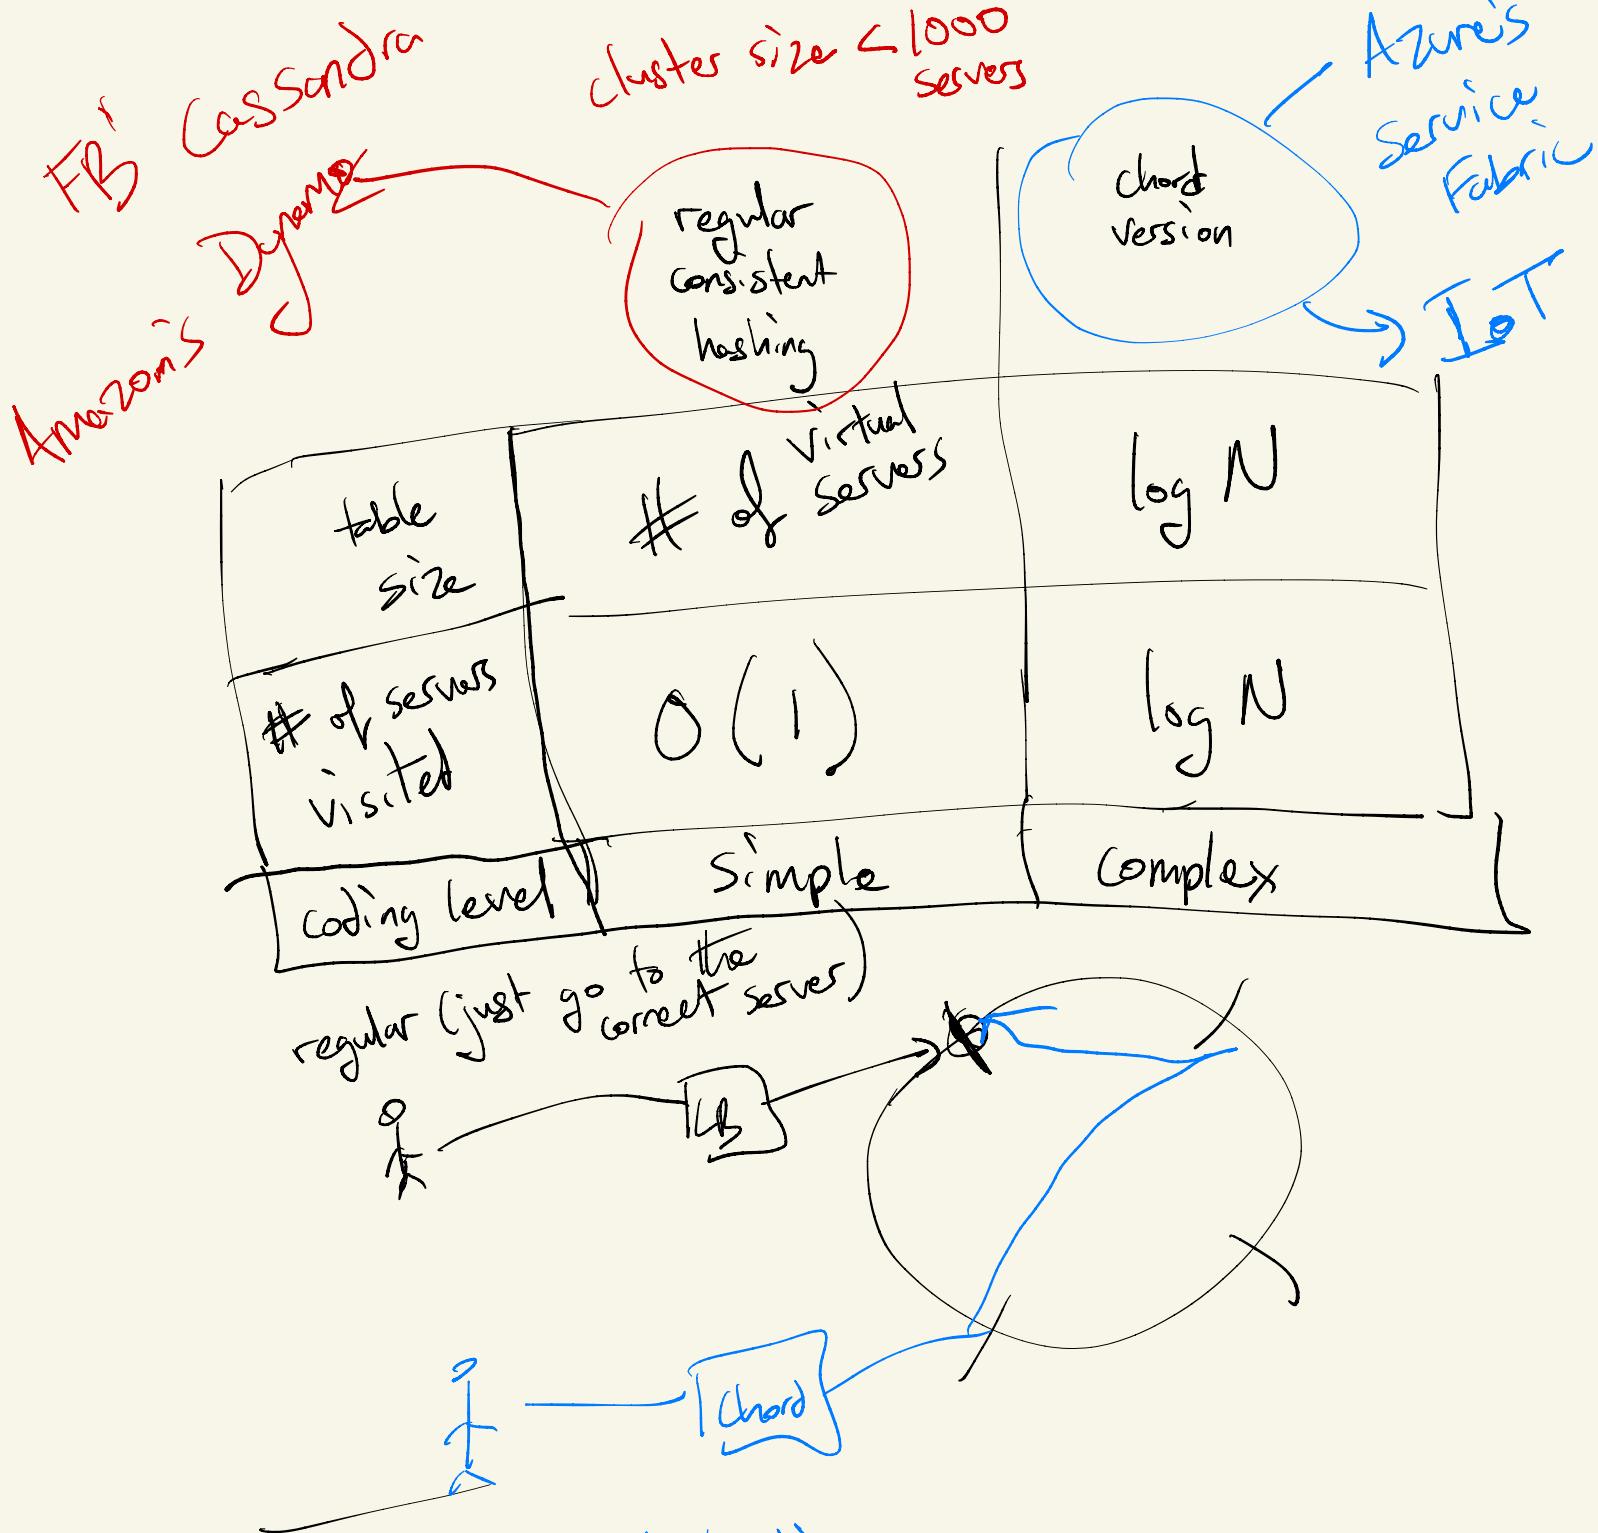
(1) Routing table size = $\log_2(N)$ key space

(2) Each entry in table is $myID + 2^{i-1}$



note my ID is
the server ID
of the current
server

(3) at each server, you determine the next
server (or the server of interest)
by checking for the largest ID smaller than
the client ID



Chord - need to binary search through several servers

$$\text{Table size} = \log_2 N = \log_2 64 = 6$$

latencies

n/w \Rightarrow 2-3 ms
CPU \Rightarrow < 1ms

n/w latencies are larger than
server latencies

thus we analyze big Oh based on
of servers

$$N = 64 \in \{0 \dots 63\}$$

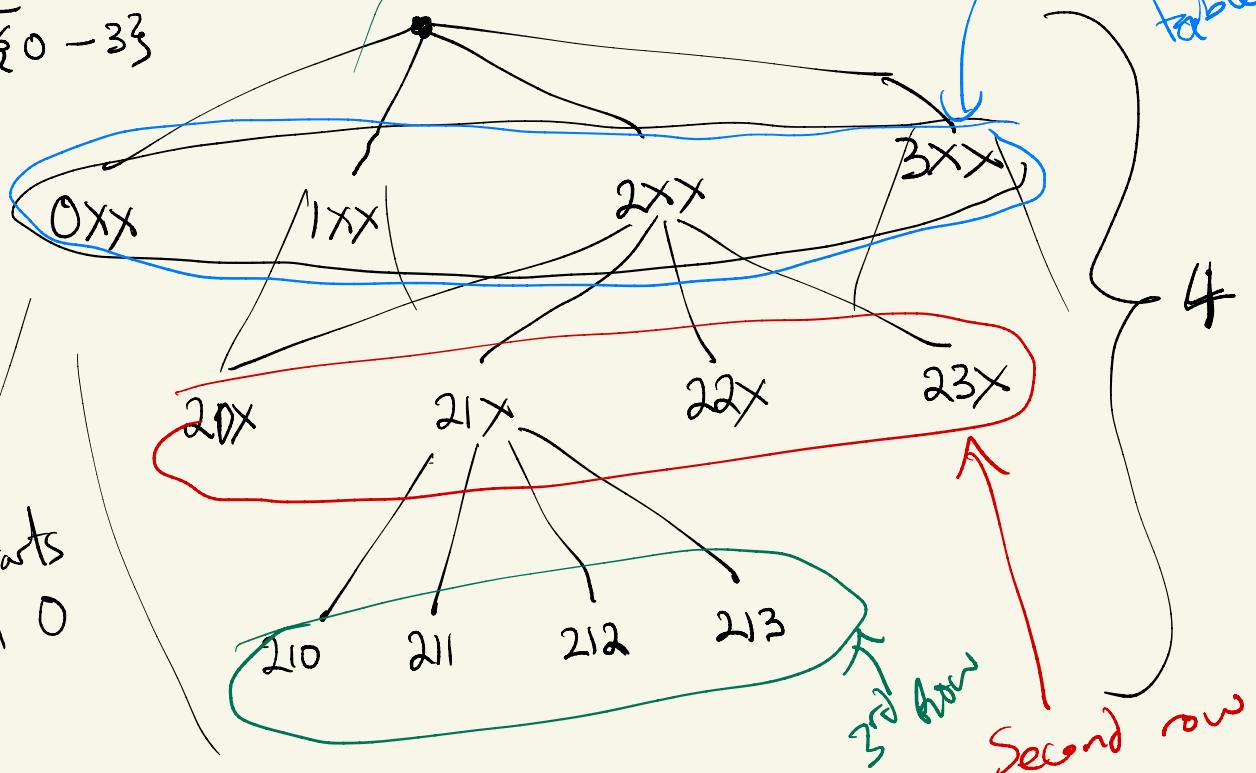
$$\text{Base} = 4 \in \{0 \dots 3\}$$

$$\log_4 64 = 3$$

$000 - 333$

Base ID Space

Starts with 0



$$\begin{aligned} N &= 64 \\ \text{Key ID Space} &= \{0 \dots 63\} \\ T &= \text{base } 10 \end{aligned}$$

Pick a Base $\Rightarrow 000 - 333$
here 4
new representation of key space with new "base"

$$\# \text{ of cols} = \text{Base}$$

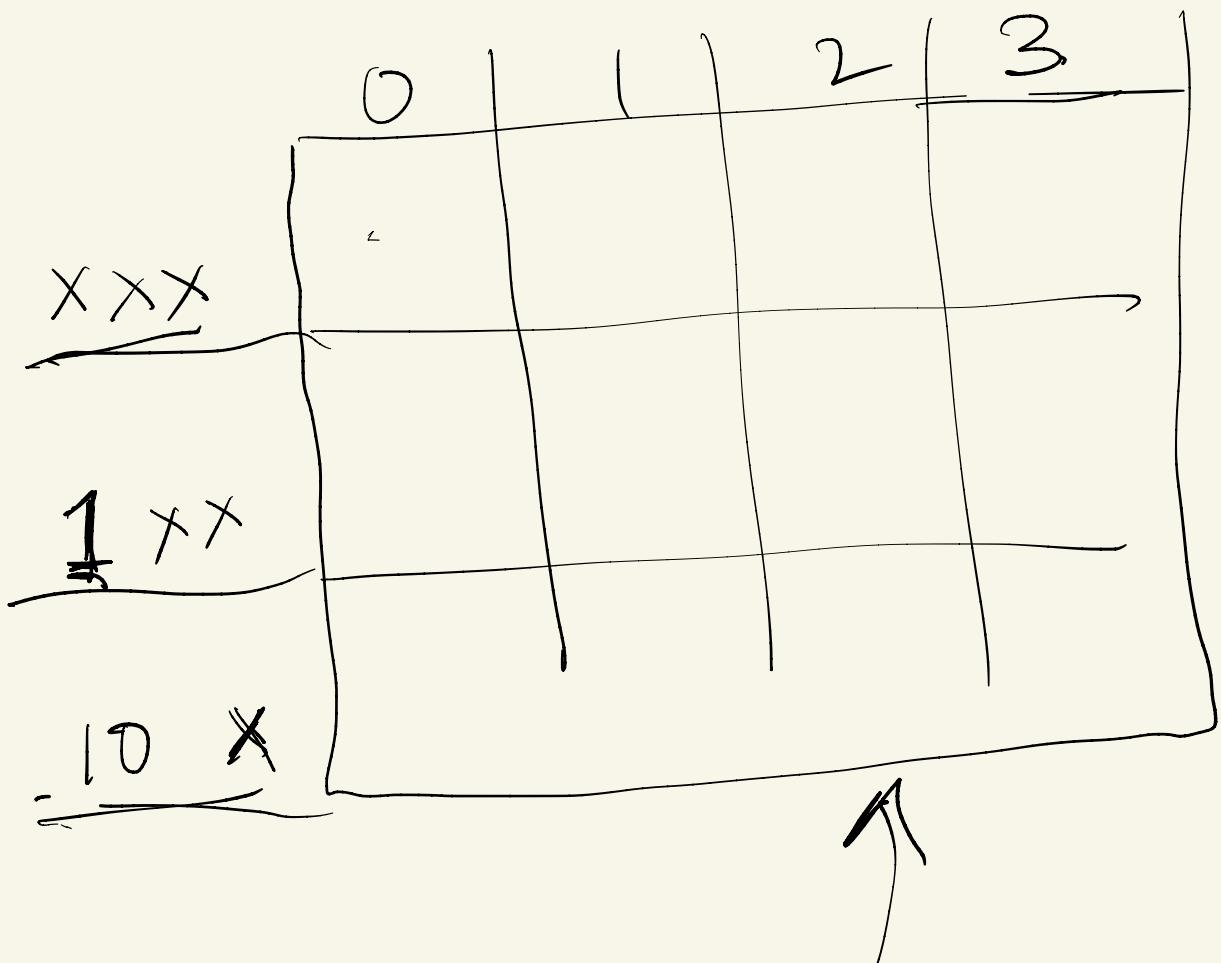
$\log_B N$			
3	0XX	1XX	2XX
	20X	21X	22X
	210	211	212

of rows

at least one node ID from each top level subtree

Conclude / Summary

- (1) assumption in consistent hash
 - (a) evenly distribution (or lack thereof)
 - (b) Virtual nodes
 - (c) maintaining a list of server ID
 - (d) chord (minimize the size of list)
(routing table)
 - (e) started a discussion on Tapestry



Server ID → 103

server ID → 321

