

CSCI 1380 : Day 10



To day

Tapestry Conclusion

Time

* why real time is bad

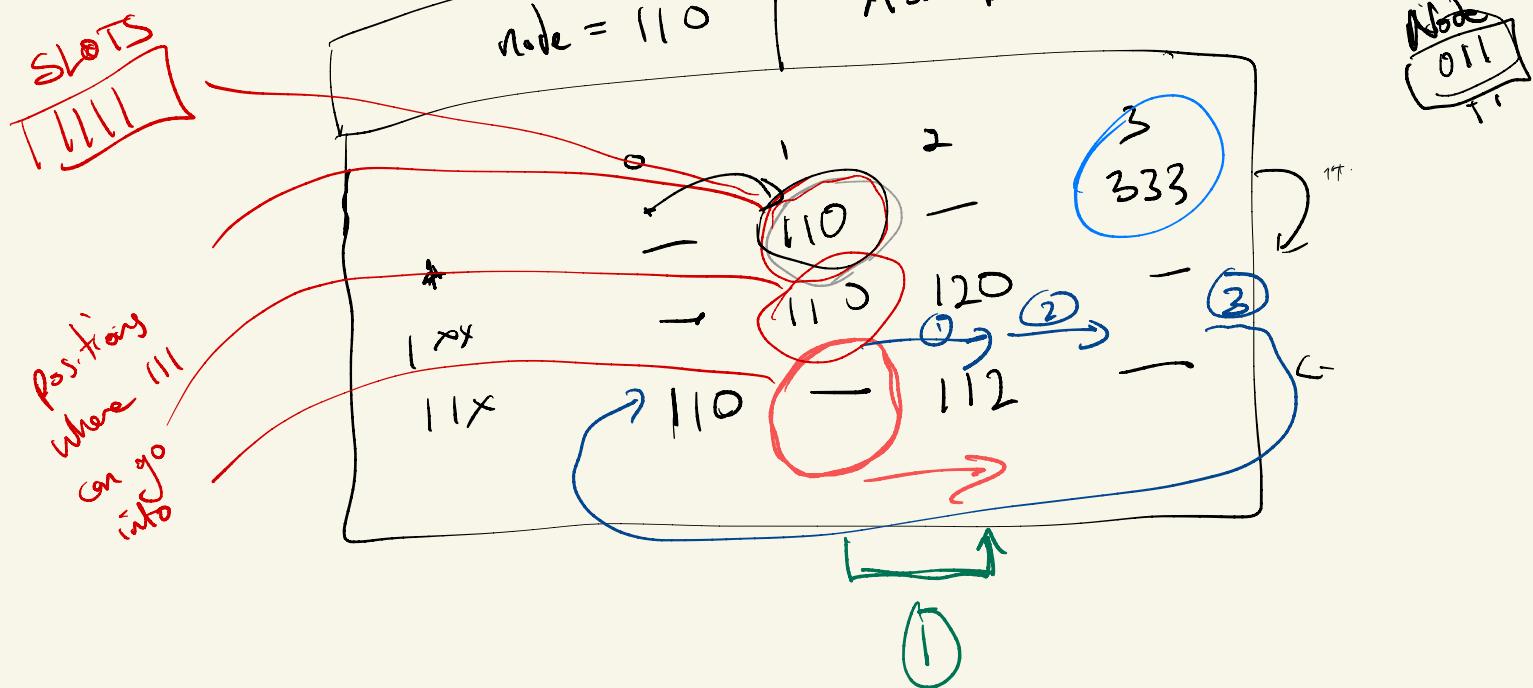
* logical clock (used in YugardB / CockroachDB)

* Vector clocks (Used in Amazon Dynamo)

Better Choice

$$(2-3) \% 4 \rightarrow (-1) \% 4 \rightarrow 3$$

Assumption : 112, 120, 333



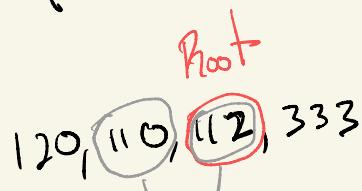
Add a new Node (III)

① find host (112)

{ ② Build Route Table

③ Inform "need-to-know" nodes
that you've joined the network

④ optimize Route table
(closes)



"need to know"

"Need to know" → a current node in
the network with a missing
cell that this new node
fills in

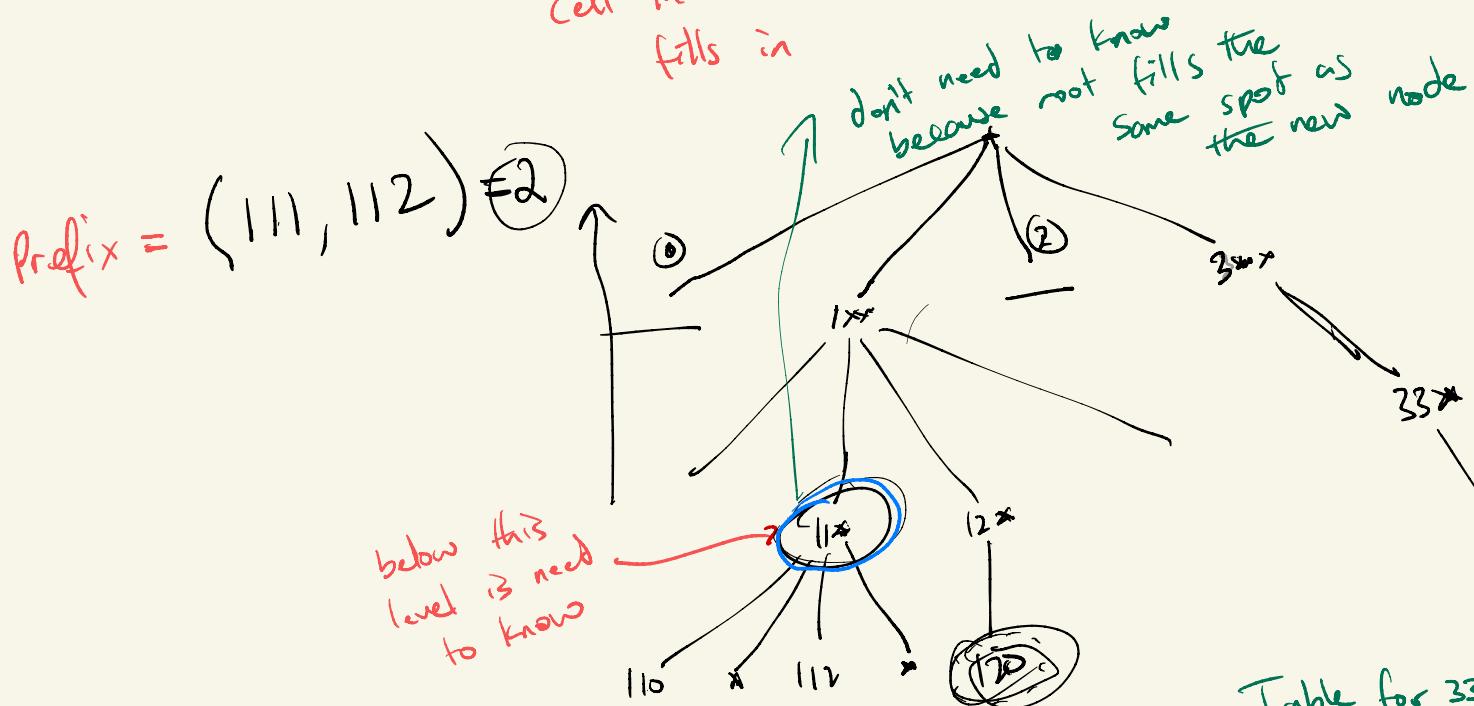


Table 120

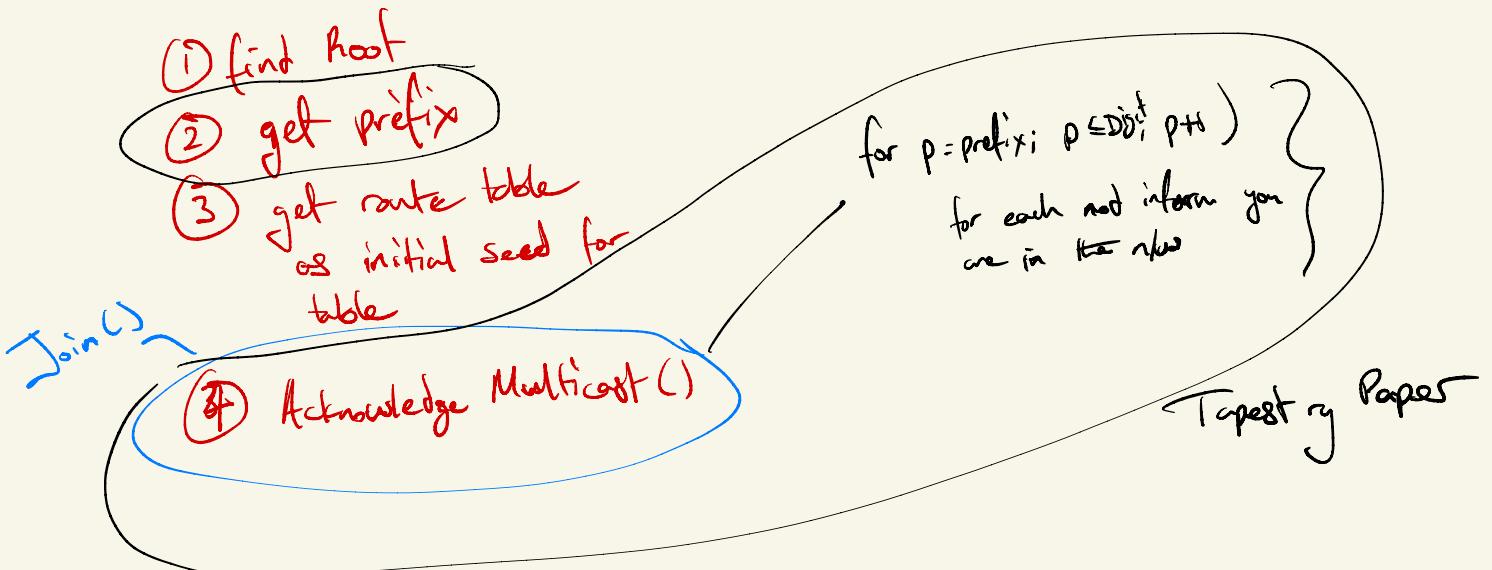
	0	1	2	3
1*				
12*	112			
*				

must start with 12*

the only place 111
can go is here but root
also meets the criteria
and would likely go
(here first)

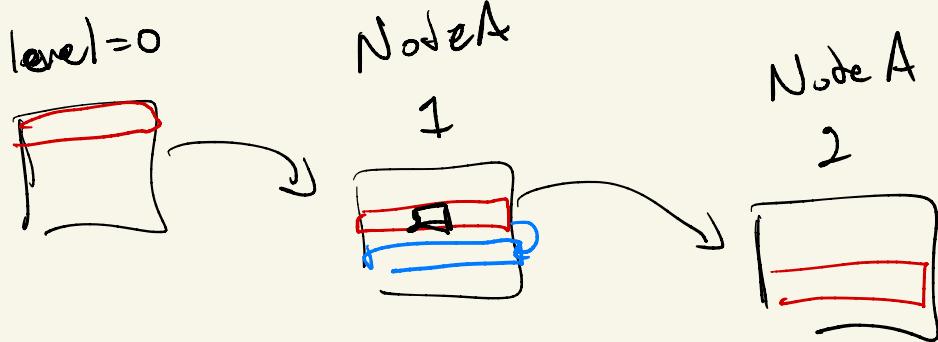
Table for 333

	0	1	2	3
3*		112		
33*				
*				



$\text{nodes} = \text{get nodes in Root's table}$
 for ($p = \text{prefix}; p \geq 0; p--$)
 $\text{BackPointer} = \text{array}[\text{in nodes}]$ {
 for each node
 $\text{backpointer.push(node, get back pointers } (p))$
 }
 $\text{local table.add(backpointer)}$.
 $\text{nodes} = \text{merge(nodes, backpointer)}$

find Neighbors



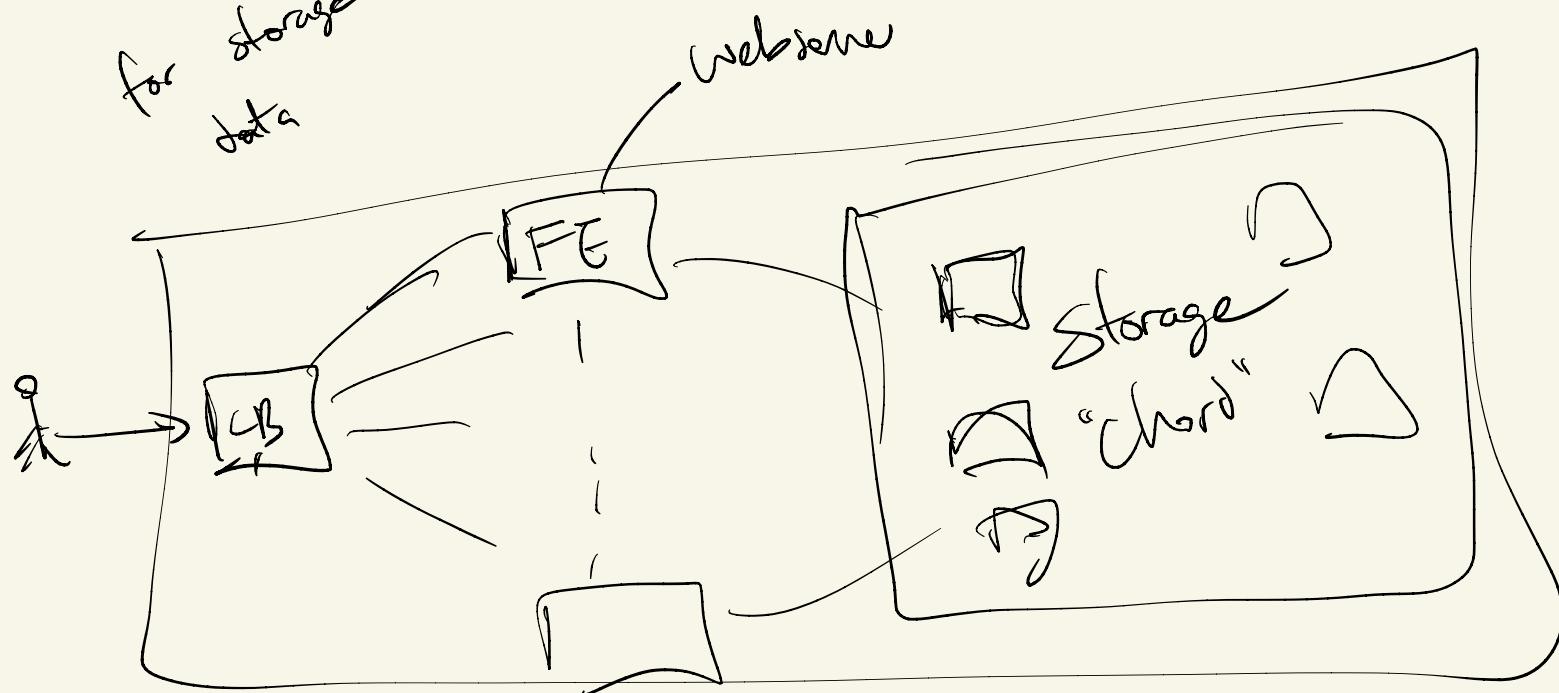
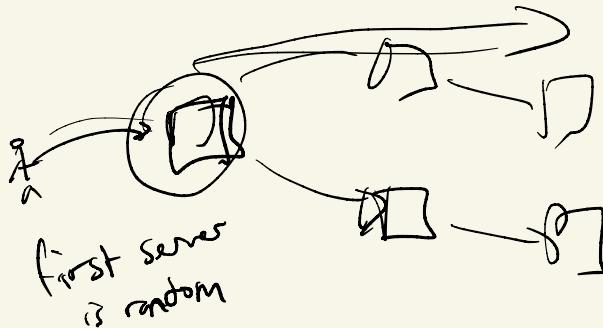
This is a recursive call from node to node where level changes between lookups

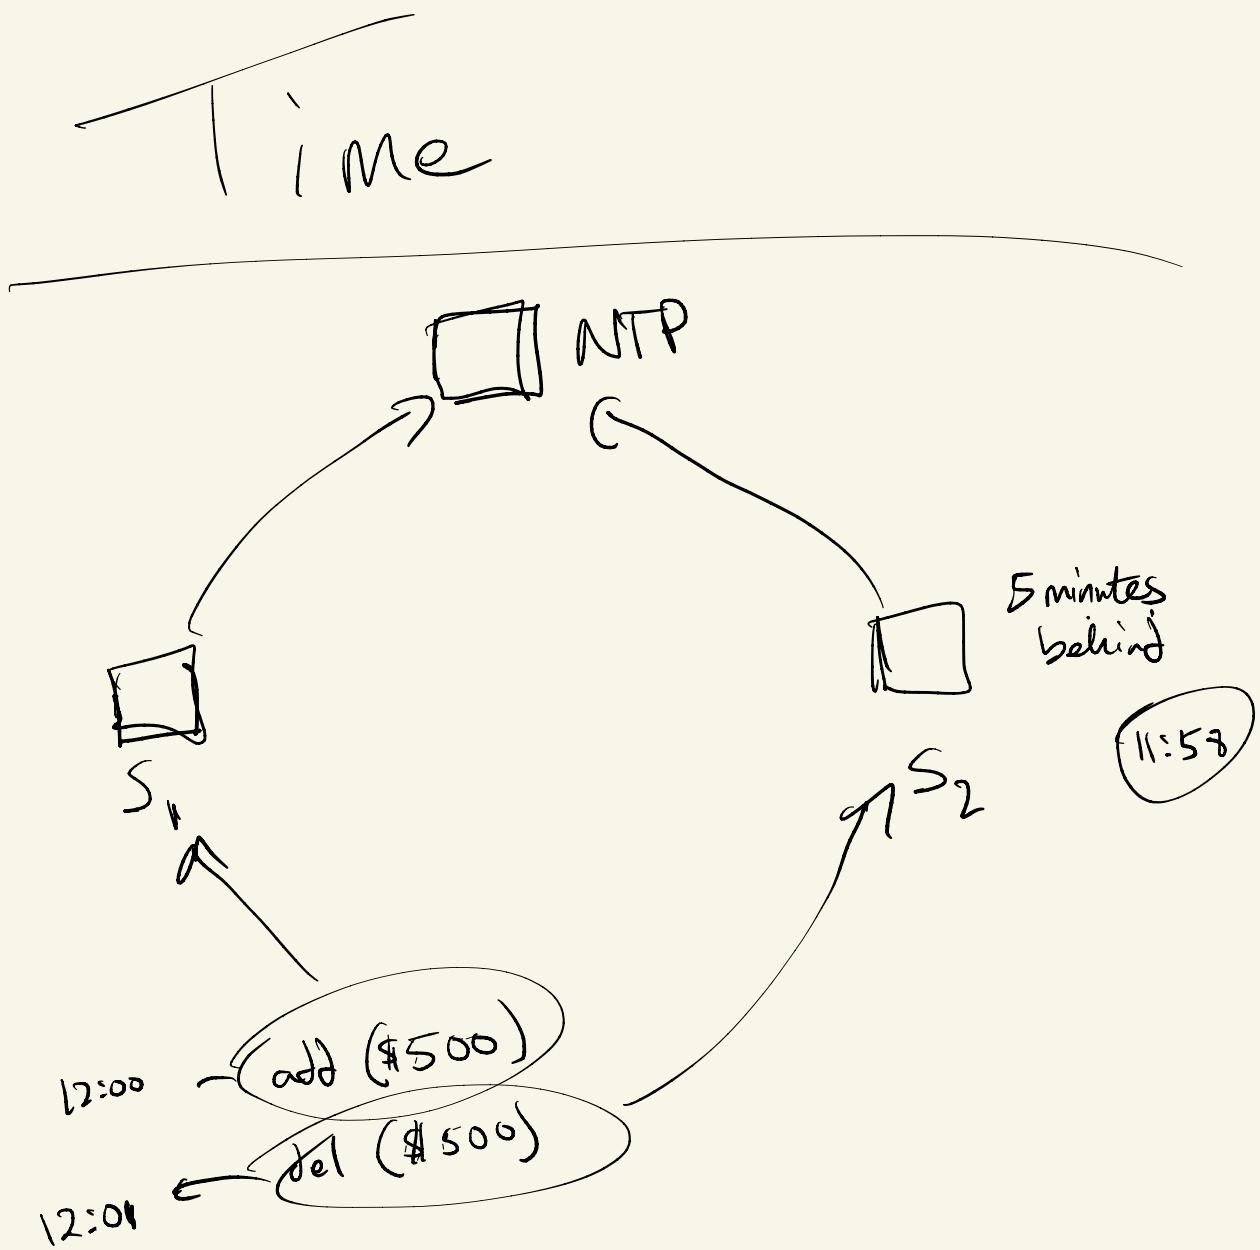
If the lookup happens at a row (say 1) of Node A but there is only one entry in the row, i.e., Node A is the only entry then you have two options

- ① recursively call Node A & go down a row
- ② simply go down a row

	Consistent hash	Chord	Tapestry
# of servers	LB	Storage	Storage
uses	IS	$\log_2(N)$	$B(\log_B(N))$
table size	$O(1)$	$O(\log_2(N))$	$(\log_B(N))$
# of servers to lookup			

Azure's Service Fabric
for storage of microservice
data





Conclusion

- ① Discussed common problems
 - a) Close
 - b) findNextHop
 - c) Join()
- ② Compared different consistent hash implementations
- ③ Started a discussion on Time