

CSCI 1380 : Day 19



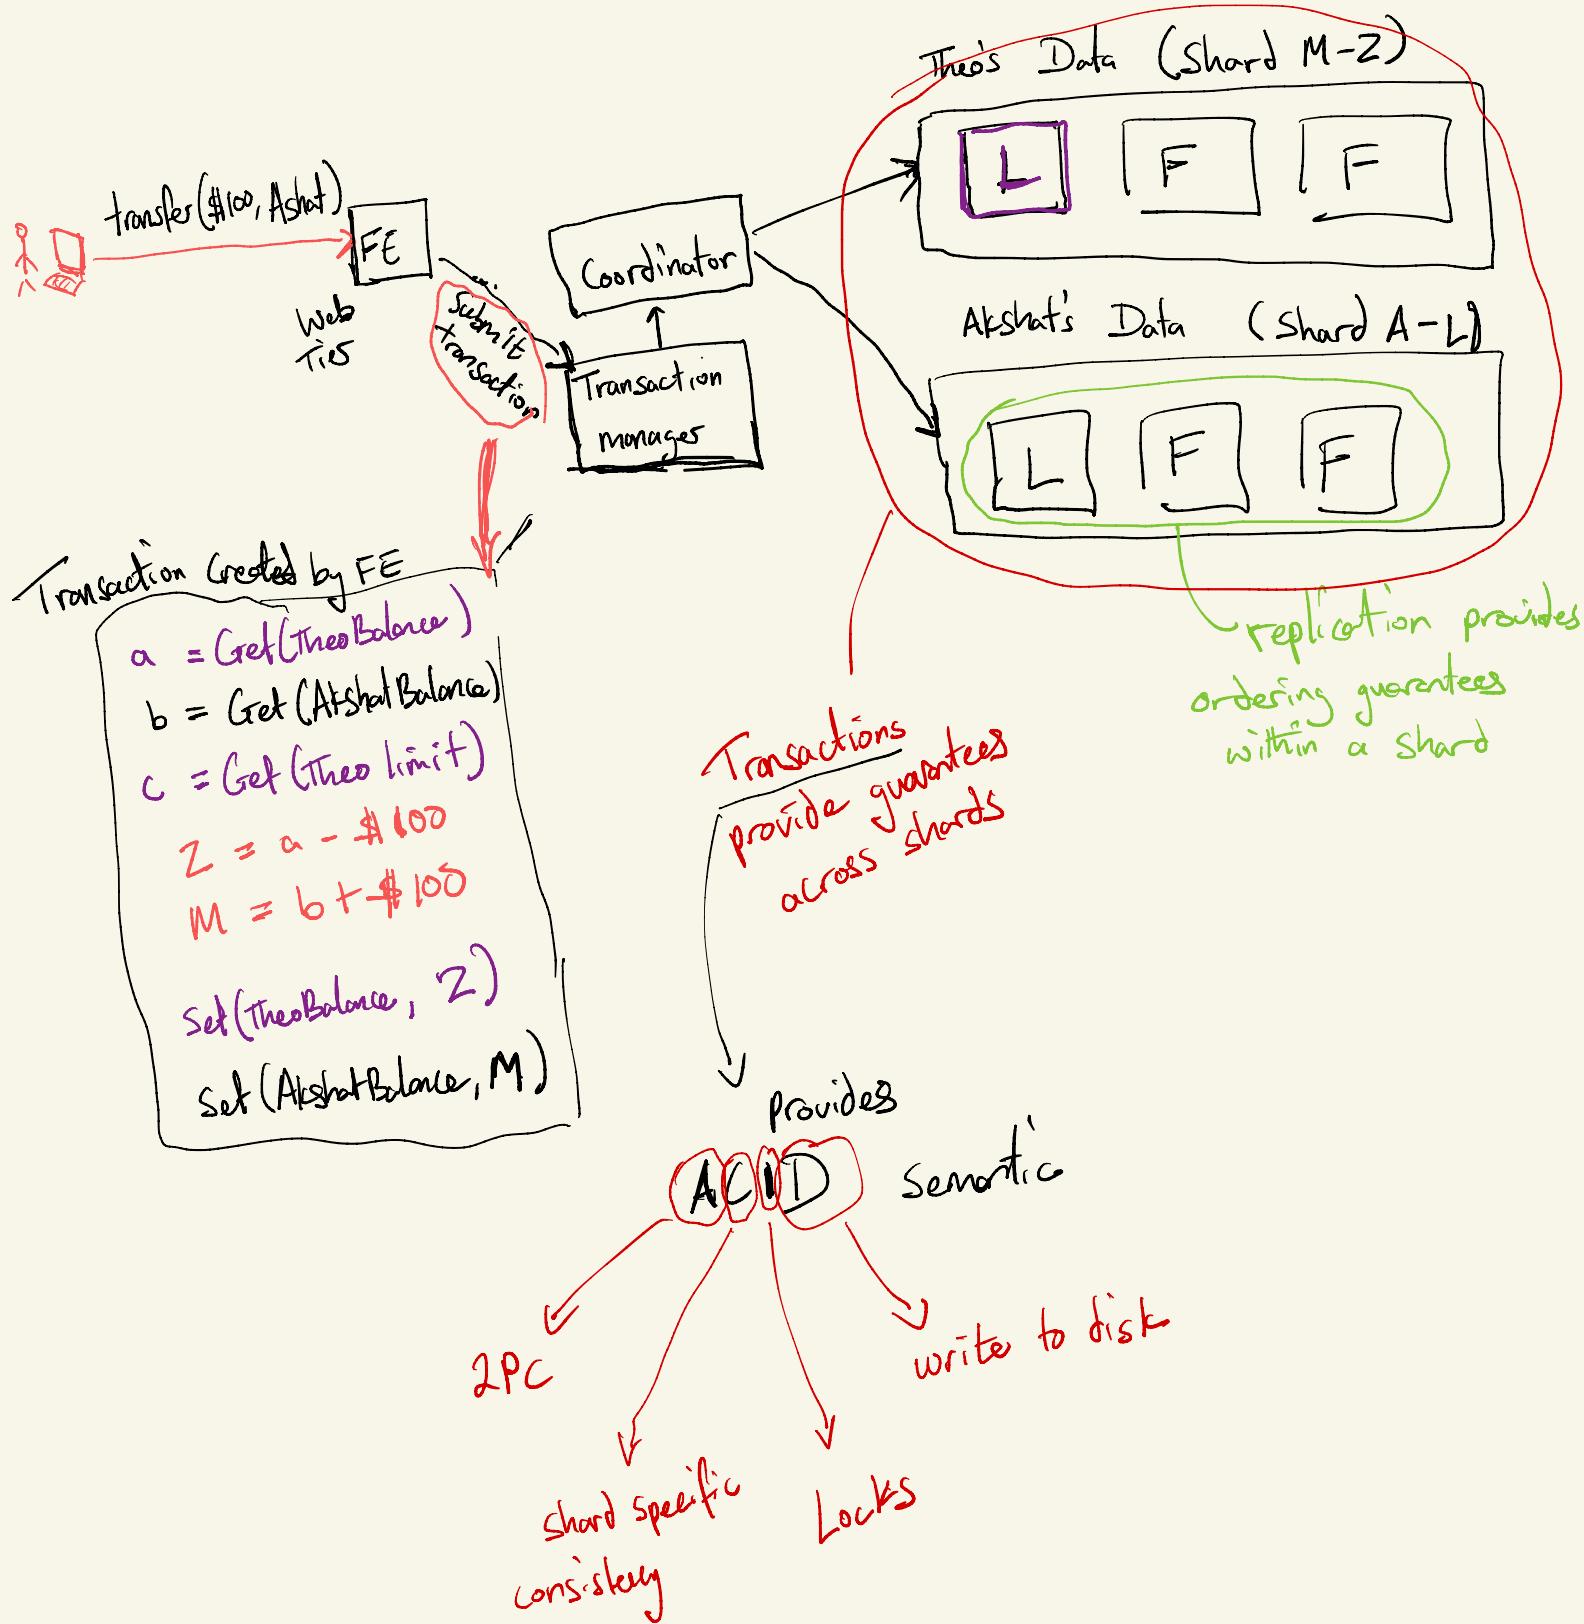
Today

Distributed Transactions concluded

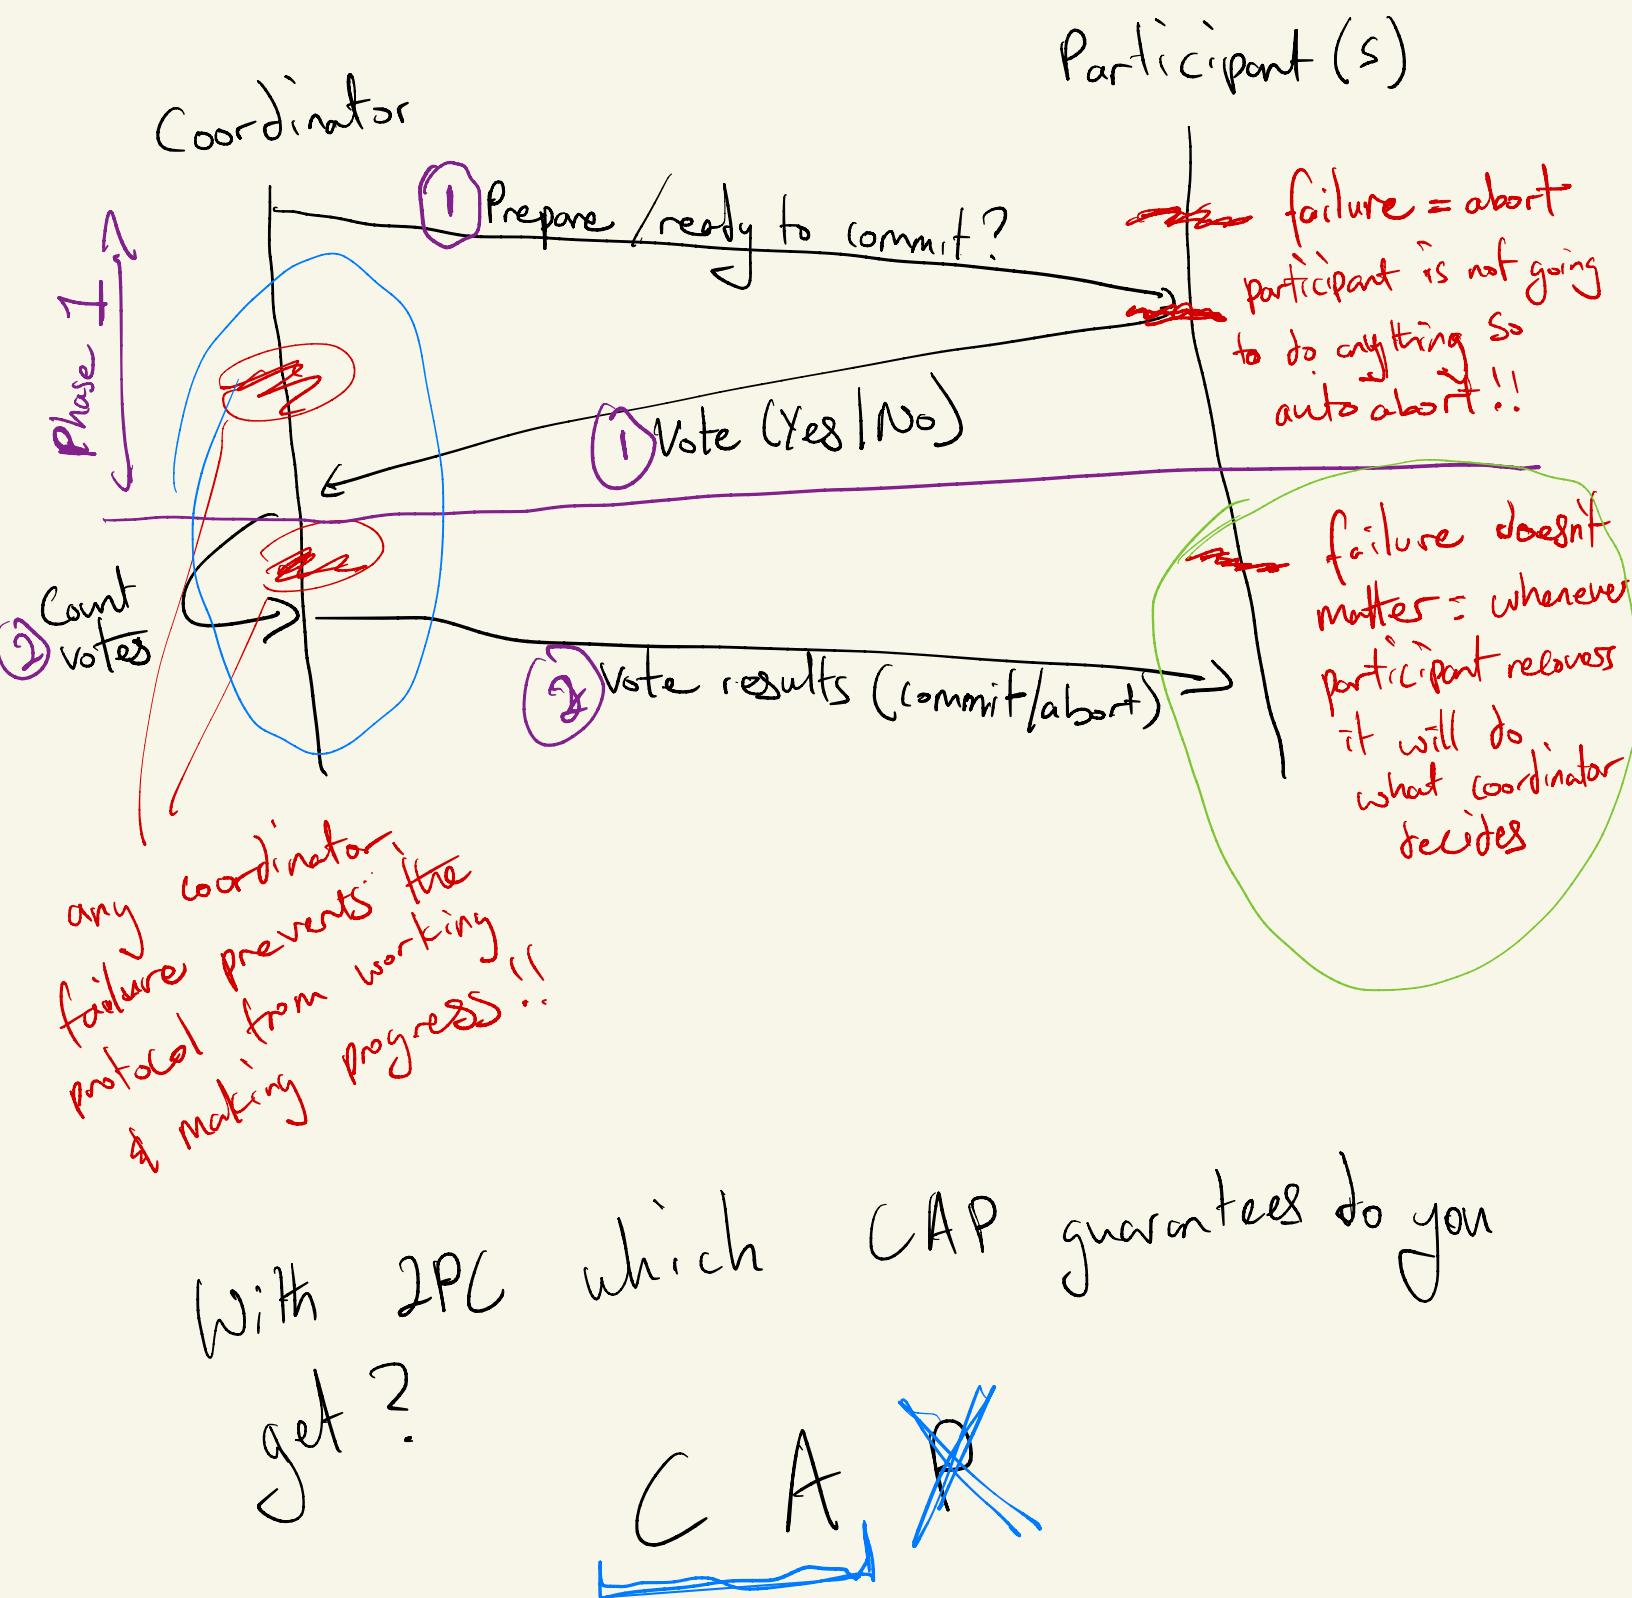
- (a) 2PC
- (b) Isolation (locking)

Distributed file systems

- (1) Scalability : metadata V. data
- (2) Performance : locks V leases & caching
- (3) Security : access control list V. capabilities

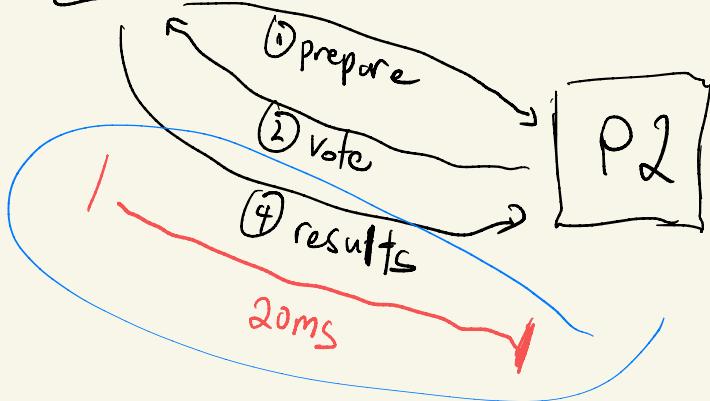
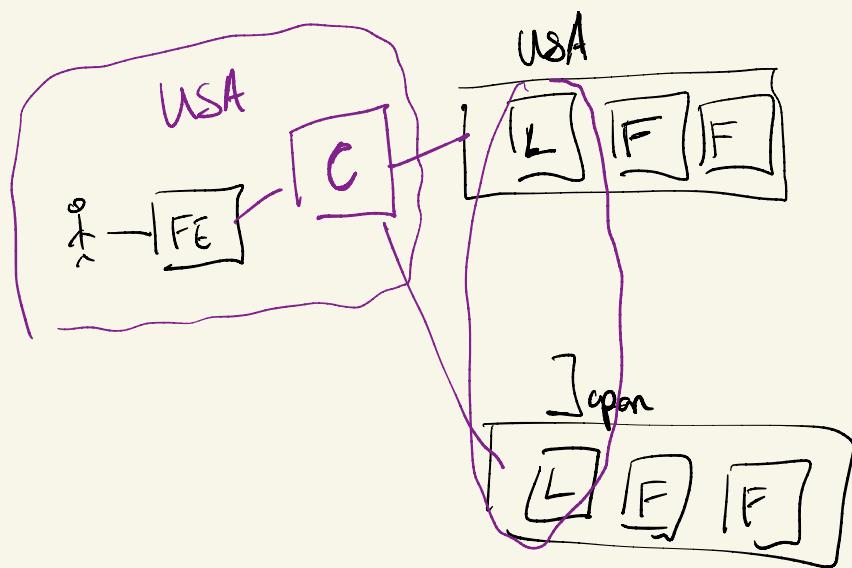
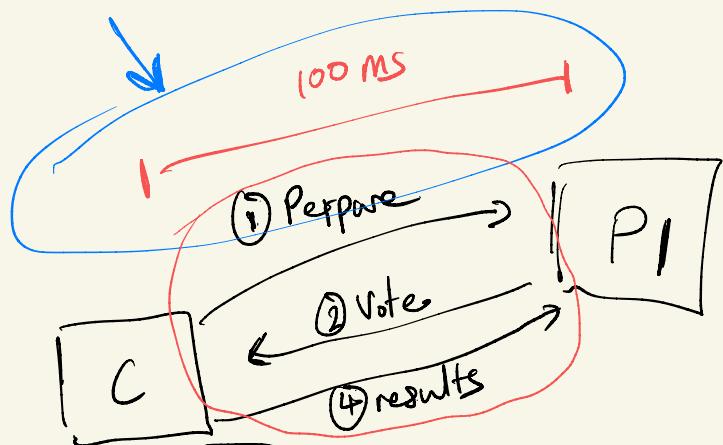


2PC

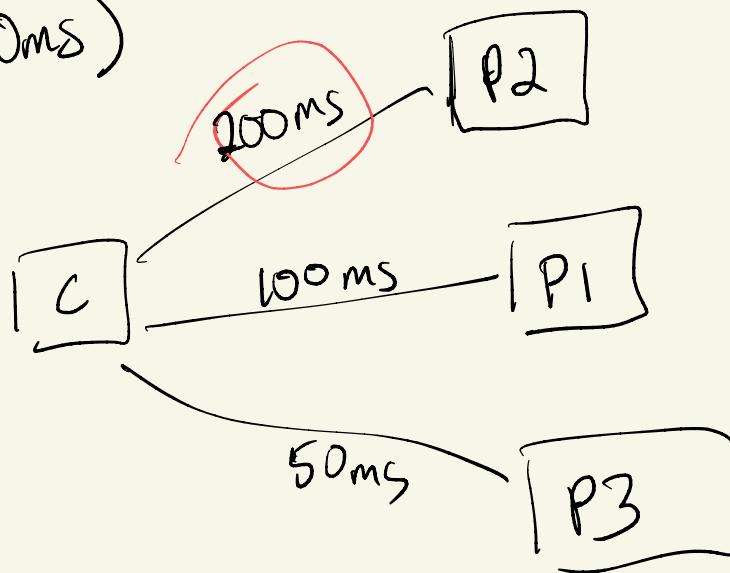


Problems with 2PC

The speed of 2PC is as fast as the slowest network latency



$$2PC = 3(200ms)$$



Providing Isolation (locks)

$a = \text{Get}(\text{TheoBalance})$
 $b = \text{Get}(\text{AkshatBalance})$
 $c = \text{Get}(\text{Theo limit})$

$$Z = a - \$100$$

$$M = b + \$100$$

$\text{Set}(\text{TheoBalance}, Z)$

$\text{Set}(\text{AkshatBalance}, M)$

Pessimistic

- ① get all locks
- ② execute all actions
- ③ commit
- ④ release all locks

Optimistic

- ① read all the data you need
- ② execute all actions
- ③ check for conflicts
- ④ if no conflicts then commit

- ⑤ rollback if conflicts

Problems

- ① Deadlocks

- ② Slow performance
(transactions execute one at a time)

Problems

- ① can have lots of rollbacks if every one needs the same data

Pessimistic = whole transaction is protected by locks

Optimistic = only a small of the transaction is protected by locks

Optimistic \Rightarrow ① perfect when transactions
use different

② transactions operate in
parallel

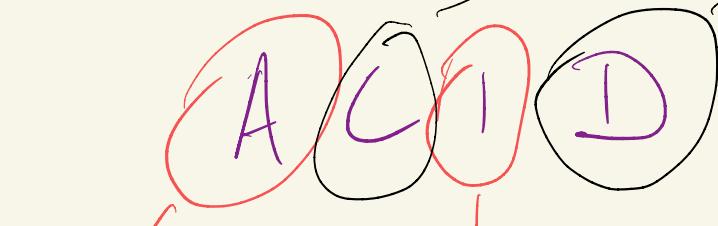
high level of
concurrency

Pessimistic \Rightarrow ① transactions operate
one at a time

② a) serial execution creates
perf. problems

b) locks can lead to
deadlocks

Dist Transactions



2PC

locking
mechanisms

intuition
for realizing
these
properties

Distributed File systems

Data

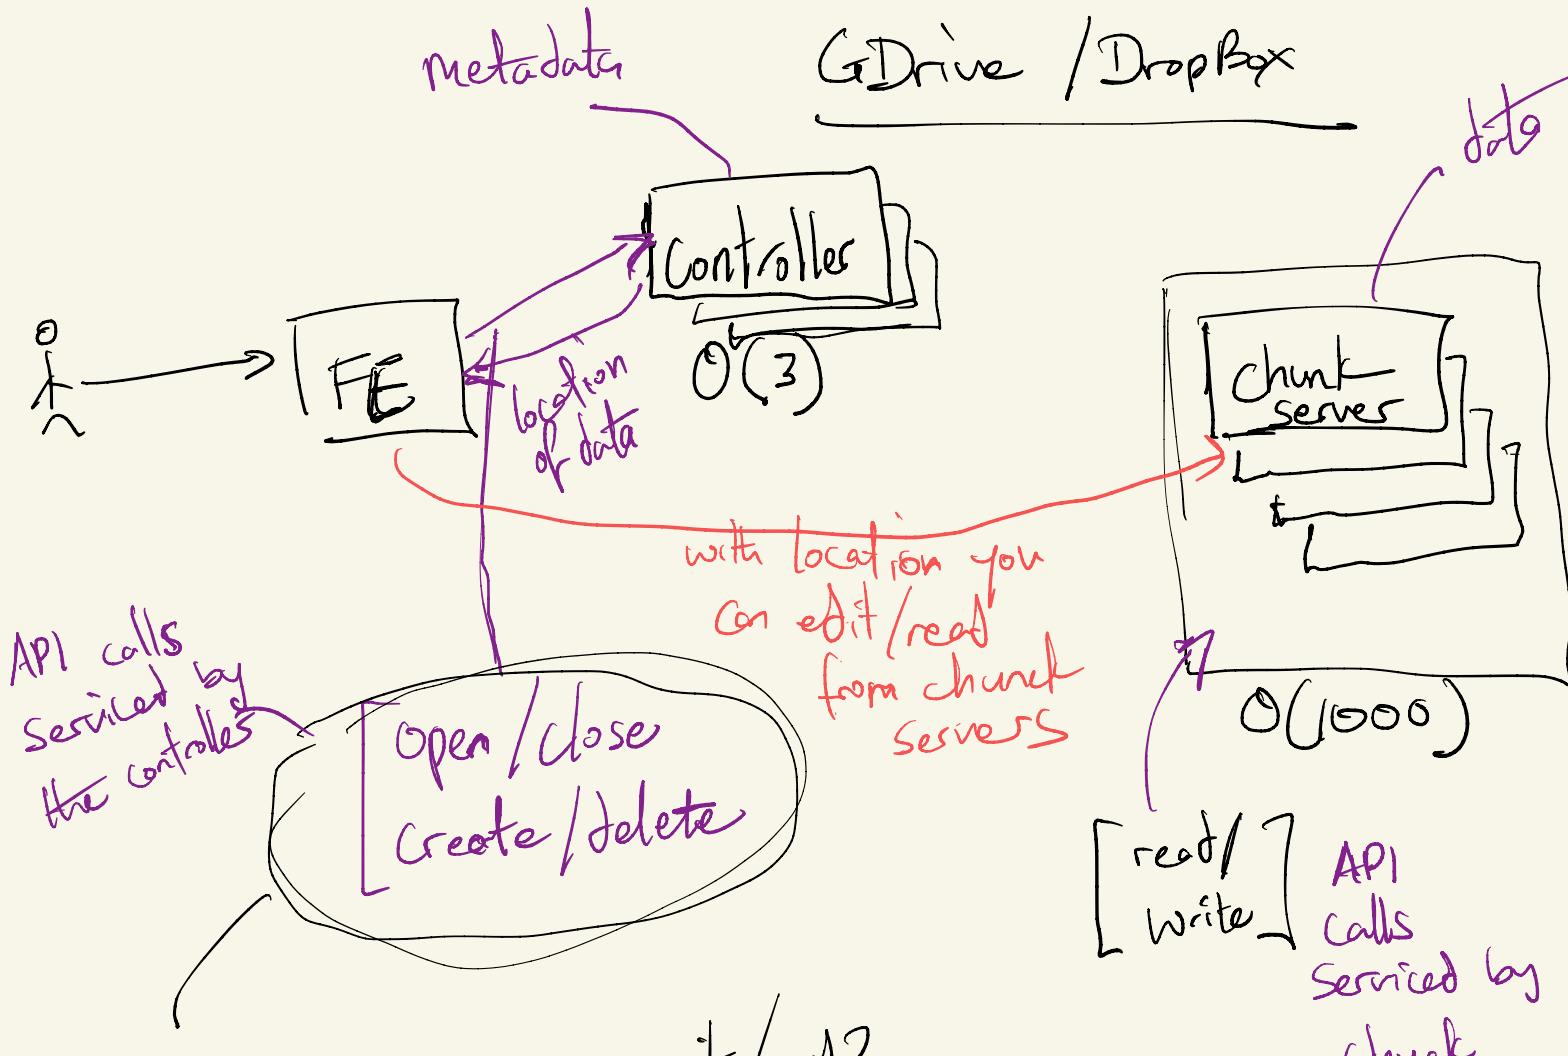
U.

meta Data

contents of the
file

is data about the data

- ① file size
- ② creation date
- ③ Access Control
- ④ File type
- ⑤ Directory



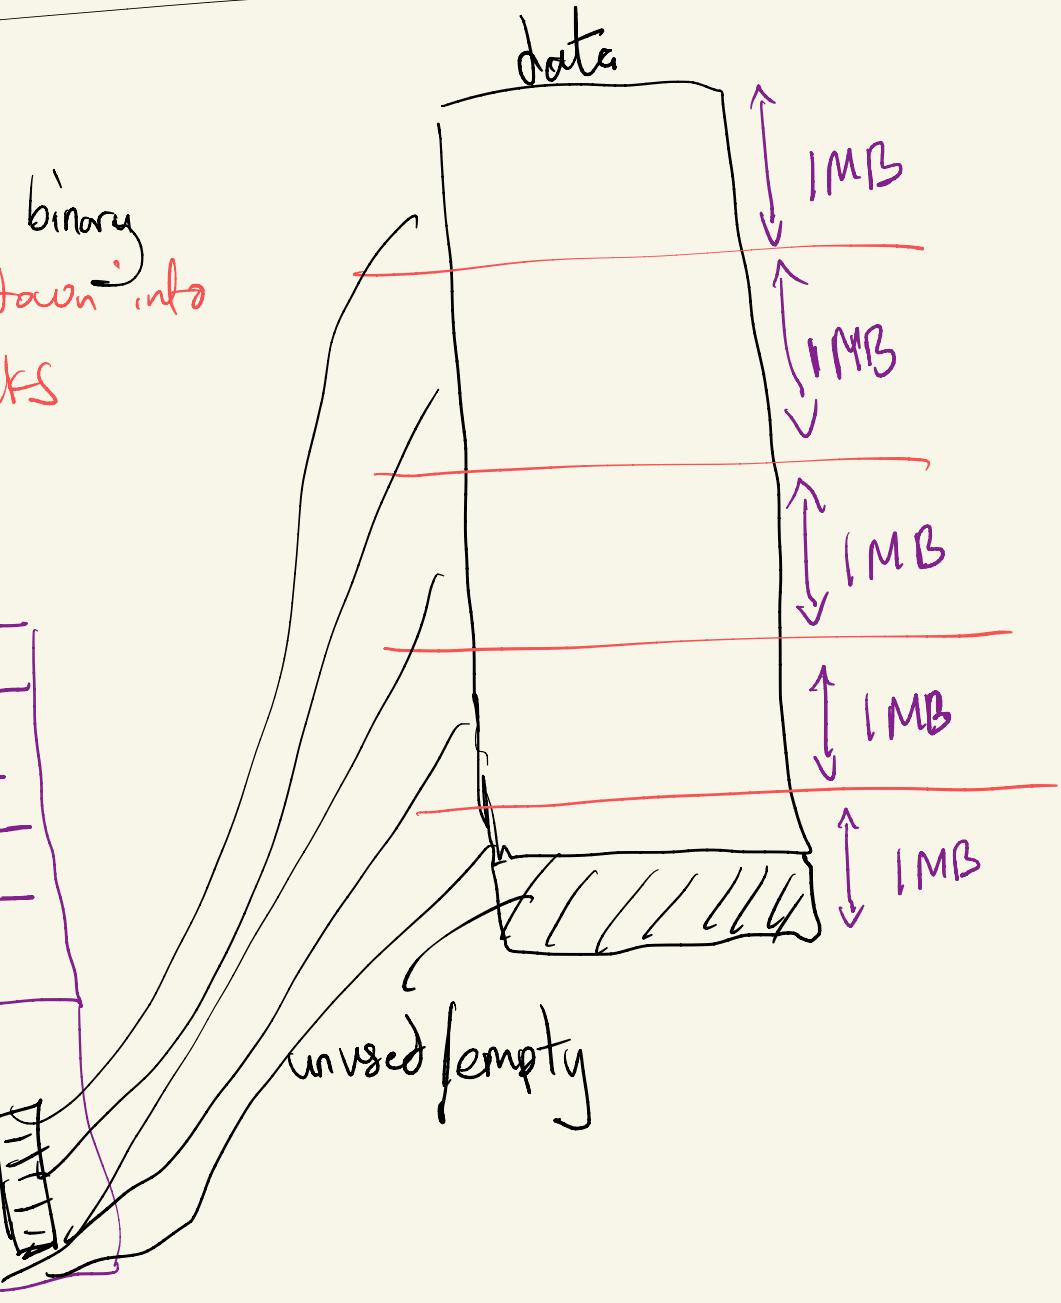
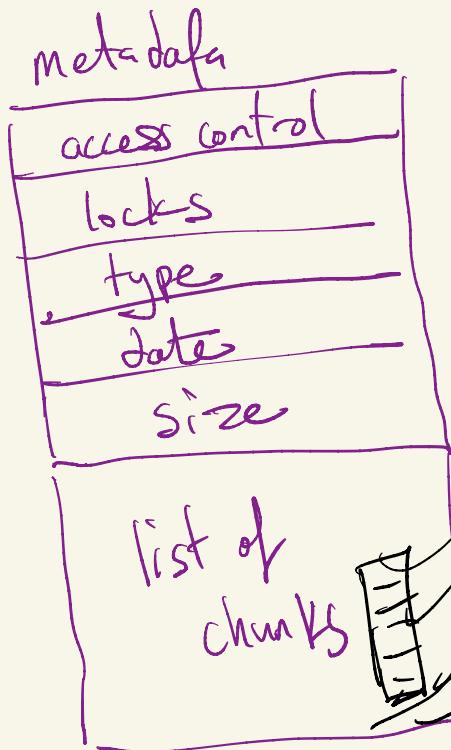
- ① permissions : can you write/read?
- ② acquire locks : gives exclusive access to chunk servers!!

DFS have a small # of controllers because they only manage metadata & thus only involved in a small # of interactions

Most of the interactions (reading/writing data) happen at the chunk servers!!

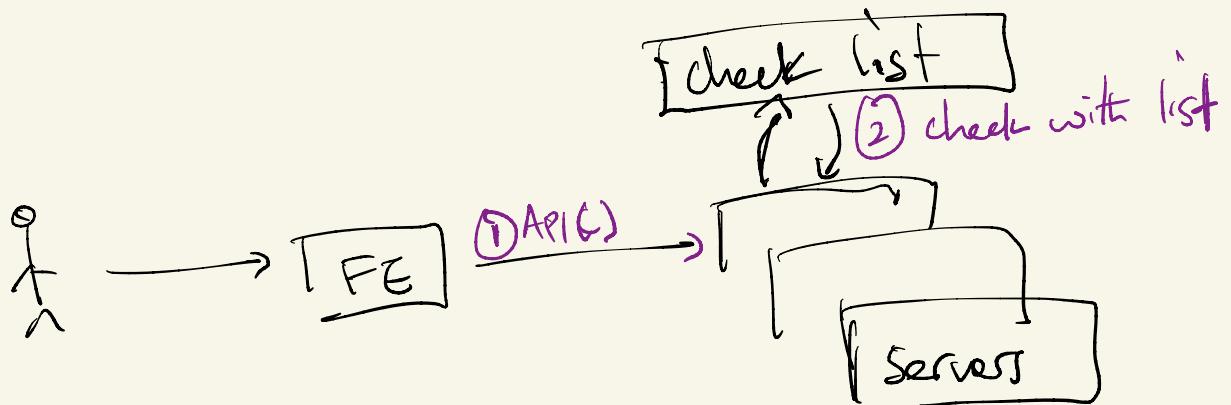
Files

File is a blob of binary
which is broken down into
fixed sized chunks



Security (Access control lists / capabilities)

Access Control list → this is a list that's maintained at servers. each API call the servers check the list to see if the client is authorized



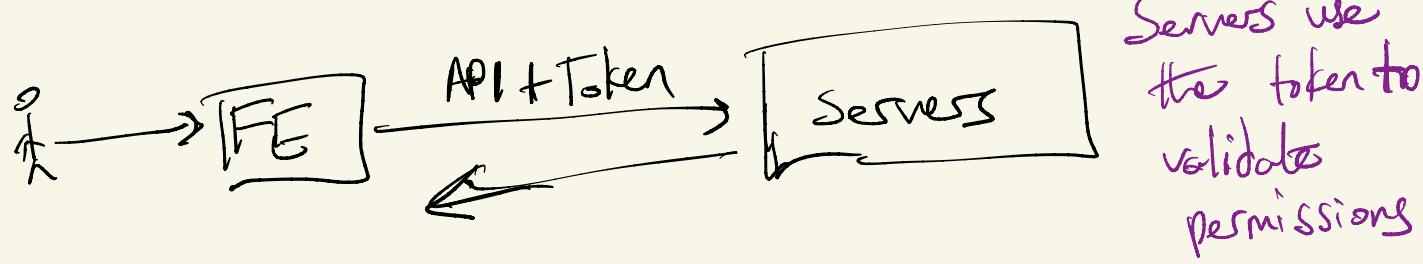
Easy to change the list & change permissions in the system

Quick to propagate security changes

Checking ACL could add significant latency for each call!!!

Capabilities

Capability is a token which encodes permissions
any one with the token can use any APIs the token
enables



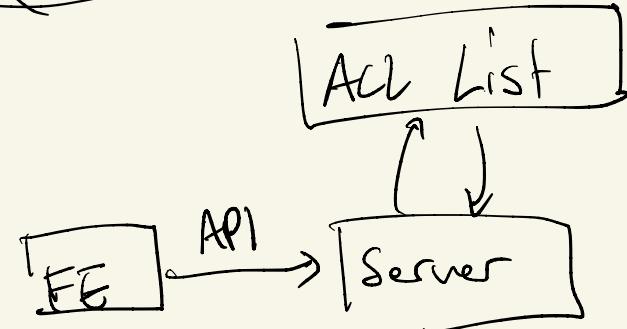
You can make copies & distribute tokens

To change permissions ; you need to track / find all tokens

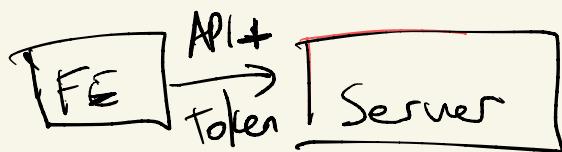
Security techniques for distributed file system

ACL

(access control lists)



Capabilities (Tokens)



① easy to permissions

② perf degradation because each call requires using the ACL

ACL

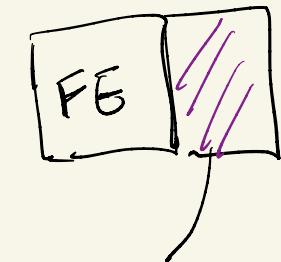
① nice & fast to verify

② Hard to change permissions because you need to track all tokens

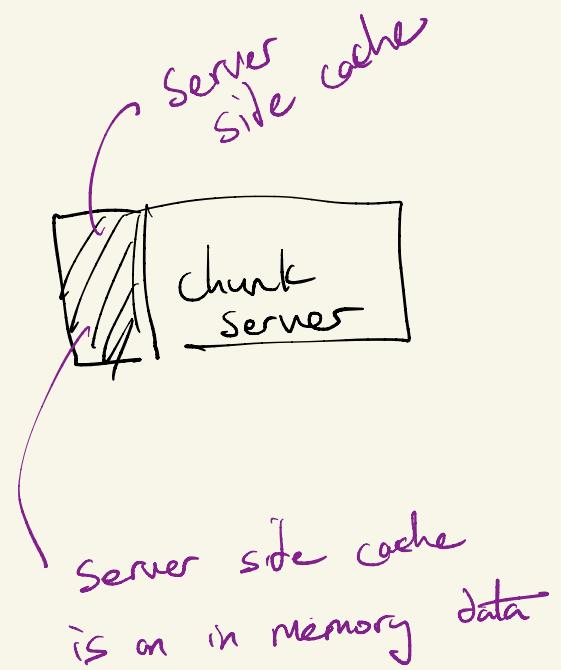
attach a timeout
to tokens

have a revocation list of bad tokens

Caching



client side
cache



① reduces load on the server (because you don't need to go to server)

The goal of this cache is to speed up reads/writes because the operations are in memory.

② all operations to cache are extremely fast because they are local (no network or server is involved !!!)

Memory >> Disk

since memory is faster than disk \Rightarrow if operations are in memory then they are faster

Today

- (1) 2PC
- (2) Isolation
- (3) Distribute file system